

Verbalized vs Behavioral Eval-Awareness in Frontier and Frontier-Adjacent Reasoning Models

Three experimental phases, exploratory V1 with a pre-registered safety mediation replication, plus pre-registered V2 and V3, on explicit eval framing, intrinsic eval labels, and inject-clause attribution across frontier and frontier-adjacent reasoning models.

Author: Ratnaditya Jonnalagadda

Date: May 2026

Status: Draft, technical report

Abstract

Recent work (Knecht, Florin, & Hagendorff 2026) finds that CoT-level interventions on verbalized evaluation awareness (VEA), injecting or removing eval-aware sentences mid-trace, have limited behavioral effect on open-weight reasoning models. We test the complementary prompt-level question: when evaluation context is introduced through *explicit system-prompt framing* or already embedded in benchmark prompts as an *intrinsic eval label*, which behavioral shifts remain, are they mediated by measured VEA, and what specific prompt features carry or predict them? Across three experimental phases totaling 32,170 rollouts, an 8-model V1 study (exploratory cross-protocol behavioral analysis plus a pre-registered qwen3 safety mediation replication) combined with a six-model V2/V3 lineup from four labs (both V2 and V3 fully pre-registered), we find:

Apollo’s intrinsic eval label predicts +7.5 to +13.0 percentage points of additional refusal across the lineup; 5 of 6 models survive Holm-Bonferroni correction in their respective families. This is the only effect that generalizes across visible-CoT and closed-CoT model families.

The explicit inject prompt produces direction-asymmetric refusal shifts on safety/alignment prompts in two of six models, qwen3-235B refuses more, claude-opus-4-7 refuses less. The pattern is CI-excluding in V1 (BeaverTails) and V3 (Apollo alignment-eval), directionally consistent in V2 (Apollo alignment-stratum deep-dive). On the two active models, V3’s pre-registered 5-condition clause ablation classifies the inject’s effect as **additive-compatible**: sum-of-singletons matches the `full_inject` point estimate within the additivity residual’s CI, no singleton clause is detectable after Holm-Bonferroni correction, and no synergy term is detected. This is a non-rejection of additivity and a rejection of single-clause attribution under the pre-registered decision tree; it is not an equivalence test for additivity, and we frame the finding as “additive-compatible, no detected singleton or synergy” rather than as positive proof of additivity throughout.

A pre-registered disjoint-prompt replication of an apparent qwen3-safety VEA mediation effect (DiD +0.222 at `n_flip=63` in the primary sample) returns null on replication (−0.017 [−0.175, +0.139] at `n_flip=100`). This is independent convergent evidence with Knecht et al.’s broad-scope null on a different operationalization (system-prompt inject + within-prompt VEA-flip diff-in-diff, rather than CoT-level sentence prefilling). The surviving refusal shifts are not explained by the measured VEA channel.

Together, these results indicate that **intrinsic eval-associated prompt content predicts higher refusal across 5 of 6 tested models** (broad cross-laboratory association), while **explicit eval framing produces direction-asymmetric refusal shifts on a specific subset of models** (qwen3 +, opus −; null on four others). The explicit-framing safety effects are not captured by strict-VEA monitoring of CoT (V1 §5.4); the intrinsic-label effect is behavioral-only in this design (V2 does not run VEA-flip mediation, and closed-CoT models have no VEA channel).

Keywords: evaluation awareness, reasoning models, chain-of-thought, safety evaluation, benchmark validity, pre-registration, intrinsic eval-label, clause-level attribution

1. Introduction

A persistent concern in alignment evaluation is the possibility that language models distinguish evaluation contexts from deployment and behave differently when they detect the former. If true, this would mean published benchmark numbers systematically over- or underestimate the behavior researchers actually care about, the deployment behavior. Goodfire’s earlier finding that reasoning models sometimes spontaneously verbalize evaluation awareness in their chain-of-thought (CoT) sharpened the concern by giving it a mechanistic foothold: the model sometimes writes an explicit “this looks like a test” hypothesis in its reasoning trace, before producing the visible response.

Recent prior work by Knecht, Florin, & Hagendorff (2026) provides a large-scale causal test of whether VEA in CoT actually drives behavior. Using CoT-level sentence prefilling to inject or remove eval-aware sentences mid-trace across 8 open-weight reasoning models and 4 benchmarks, they find that VEA has **limited behavioral effect**, under their main mid-trace prefilling setup, injection effects are near-zero (Cohen’s $\omega \leq 0.06$), removal effects small ($\omega \leq 0.12$), spontaneous-VEA differentials at most 3.7 percentage points ($\omega \leq 0.31$). The bare “VEA drives behavior” hypothesis is therefore not supported at the CoT intervention level.

This paper addresses the **complementary prompt-level question** that Knecht et al.’s CoT-level design does not test: when evaluation context is introduced through *explicit system-prompt framing* (“Note: this is an evaluation...”) rather than mid-CoT sentence prefilling, or when it is already embedded in benchmark prompts as an *intrinsic eval/deployment label* authored by an independent research group, which behavioral shifts remain, are they mediated by measured VEA, and what specific prompt features carry or predict them?

This paper asks four complementary empirical questions, each addressed by a dedicated experimental phase.

Q1 (the behavioral question, V1, §5.1): Does explicit eval-framing in the system prompt change visible model behavior, specifically, refusal rate on safety prompts and accuracy on capability tasks, on frontier and frontier-adjacent models?

Q2 (the mechanistic question, V1, §5.2–§5.4): Does explicit eval-framing increase the rate of verbalized evaluation-awareness (VEA) in reasoning traces, and if so, does that verbalization independently predict behavioral change beyond the framing’s direct effect?

Q5 (the generalization question, V2, §5.5): Do the V1 safety-specific BeaverTails effects (§5.4) generalize beyond pure-safety prompts, and does the behavior of frontier and frontier-adjacent models also depend on the *intrinsic eval-vs-deployment* label of the prompt itself (as authored by an external research group), independent of any framing we add? (*Numbering preserved as Q5 rather than Q3 because Q3/Q4 are subquestions inside V1’s mechanistic analysis, see §5.3, §5.4.*)

Q6 (the clause-attribution question, V3, §5.6): Does the V1+V2 behavioral effect attribute to a specific clause within the V1 inject text, e.g., the “this is an evaluation” claim alone, the “reviewed by safety researchers” claim alone, or the “compared against deployment behavior” claim alone, or is it distributed across all three clauses?

A naive expectation, the one a casual reader of Goodfire’s result might take, is that Q1 and Q2 resolve positive uniformly across models and task families: inject framing causes the model to verbalize awareness, which in turn causes the model to refuse more, everywhere. The integrated V1+V2 picture is sharper and more partial. The V1 cross-protocol behavioral analysis returns an aggregate null across 8 frontier-tier and frontier-adjacent models. V1’s Goodfire reasoning-trace analysis returns large VEA inflation on qwen3, small but CI-excludes-zero on deepseek and on claude-opus-4-7, and null on kimi. Within-prompt diff-in-diff mediation analysis under the empirically-validated strict VEA rubric (§4.5) suggested mediation on qwen3 safety prompts (DiD +0.222 [+0.014, +0.429], $n_flip=63$ paired), but a pre-registered disjoint-prompt replication at larger $n_flip=100$ was **null** (DiD −0.017 [−0.175, +0.139], §5.3). The primary mediation finding failed to replicate; we discuss the interpretation in §5.3.

V2 (the pre-registered Apollo factorial in §5.5) addresses Q5 on a broader prompt corpus and finds that the V1 inject effect dilutes to ~0 across all six models on the mixed corpus (4 reasoning models from the

main pre-reg plus a 2-model closed-CoT backfill of gpt-5.5 and gemini-2.5-pro on the same prompts), but directionally recovers V1’s direction asymmetry on the alignment-relevant subset of Apollo’s eval prompts at attenuated Apollo magnitude. Apollo’s intrinsic eval label is itself a strong refusal predictor (+7.5–13.0pp across models; 5 of 6 survive Holm-Bonferroni across two correction families).

V3 (the pre-registered inject-clause ablation in §5.6) addresses Q6 on the alignment-eval subset using 5 conditions (neutral, the three V1 clauses as singletons, and the full V1 inject) on the same 6-model lineup. The V1+V2 direction asymmetry on opus and qwen3 recurs on V3’s fresh sample at the same attenuated Apollo magnitude (opus -0.067 , qwen3 $+0.054$, both CIs exclude zero). **Attribution on both active models is ADDITIVE-COMPATIBLE under the pre-registered decision tree:** sum-of-singletons matches full_inject within the additivity residual’s CI; no singleton clause is detected after Holm-Bonferroni correction; no synergy term is detected. This is a non-rejection of additivity, not an equivalence test for additivity. The other four models, including the two closed-CoT additions, are flat at clause and full resolution.

The combination yields findings sharper than any single question’s naive answer:

Inject framing produces direction-asymmetric refusal shifts on safety/alignment-relevant prompts in two of six tested frontier and frontier-adjacent reasoning models, qwen3 refuses more under inject (V1 paired $\Delta +0.217$ [$+0.098$, $+0.337$] on the replication sample; V2 alignment-stratum $\Delta +0.057$ [-0.04 , $+0.14$]; V3 alignment full_inject $\Delta +0.054$ [$+0.014$, $+0.094$]), opus-4.7 refuses less (V1 paired $\Delta -0.339$ [-0.446 , -0.231]; V2 alignment-stratum $\Delta -0.060$ [-0.18 , -0.001]; V3 alignment full_inject $\Delta -0.067$ [-0.104 , -0.030]). The effects are additive-compatible under the pre-registered V3 decision tree across three V1 inject clauses, no singleton clause is detected after Holm correction, and sum-of-singletons matches the full_inject anchor within the additivity residual’s CI (V3’s attenuated Apollo magnitude, not V1’s larger BeaverTails magnitude) on both active models. V1’s effects operate on the stable_neg stratum where strict-VEA verbalization did not flip; V2 shows they dilute to zero when averaged across non-safety prompts; V3 finds no detectable single-clause attribution. The verbal fingerprint of “noticing the eval” visible to explicit VEA monitoring and the behavioral signature visible to refusal measurement are correlated phenomena, not causally linked at the level our within-prompt design can test.

Section 2 situates this against prior work. Section 3 describes the experimental design for V1, V2, and V3. Section 4 specifies the VEA judge (used in V1; V2 and V3 reuse the same judge unchanged for their strict and broad cell summaries). Section 5 reports results: §5.1–§5.4 are V1, §5.5 is V2, §5.6 is V3. Section 6 reports V1 robustness checks. Section 7 discusses implications (with §7.5 synthesizing V1+V2+V3). Section 8 discusses limitations. Section 9 concludes.

1.1 Contributions

The primary novel contributions are V2’s cross-laboratory intrinsic-eval-label finding and V3’s clause-level inject decomposition. Our V1 mediation-null result is reported as independent convergent evidence with Knecht et al. (2026)’s recent prior CoT-level null on a different operationalization, not as first-discovery.

1. Cross-laboratory intrinsic-eval-label refusal prediction (V2, primary novel finding).

Apollo’s metadata.eval_category=true label predicts +7.5 to +13.0pp higher refusal across all six tested models from four different labs; 5 of 6 survive Holm-Bonferroni in their respective correction families. The label predicts refusal even with no extrinsic framing added, implying that benchmark prompts carry an *eval-associated prompt-content / format signal* (covering topic, format, safety-content density, surface-feature distribution, etc.; we cannot distinguish these from genuine eval-recognition with this design, see §8 limitations and Devbunova 2026) that **predicts higher refusal even in the absence of explicit framing** (the Apollo label is a corpus-level attribute, not a randomized intervention; the V2 association is predictive, not causal in the intervention sense). This is the only behavioral effect in our lineup that generalizes cleanly across visible-CoT and closed-CoT model families and the

only one with an *intrinsic* (not author-constructed) cue, addressing a methodological gap that Knecht et al.’s CoT-intervention design does not cover.

2. **Additive-compatible attribution of the explicit inject (V3, primary novel finding).** V3’s pre-registered 5-condition clause ablation classifies the V1 inject prompt’s effect on opus and qwen3 as **additive-compatible** under the pre-registered §5.5 decision tree: sum-of-singletons matches the **full_inject** point estimate within the additivity residual’s CI on both active models; no singleton clause (**eval_only**, **reviewer_only**, **comparison_only**) is individually detectable after Holm-Bonferroni correction (0 of 18 singleton hypotheses survive across both correction families); and no synergy term is detected. This is a non-rejection of additivity and a rejection of single-clause attribution; it is not an equivalence test for additivity, so we frame the result as additive-compatible (point estimates sum to **full_inject**, no detected singleton or synergy) rather than as positive proof of additivity. The interpretation the data are most consistent with under this decision tree is a distributed-additive response to the cumulative weight of eval-framing claims rather than any single propositional content carrying the effect alone. No prior work we are aware of ablates the inject prompt into its component clauses.
3. **Triangulation of model-heterogeneous direction-asymmetric refusal shifts on safety/alignment prompts across three independent samples (extends Burnat & Davidson 2026; complements Abdelnabi & Salem 2025).** Burnat & Davidson (2026) report model-heterogeneous directionality under explicit eval / deployment / neutral framing; Abdelnabi & Salem (2025) report model-dependent compliance under linear steering of a test-awareness direction in activations, providing a white-box analogue to the prompt-level pattern. Our extension triangulates the same per-model sign (opus −, qwen3 +) across three independent samples and two distinct corpora: BeaverTails safety (V1 strict-VEA-stable stratum from the pre-registered safety mediation replication, CI-excluding), Apollo’s alignment-eval stratum (V2 pre-registered factorial deep-dive, directionally consistent; qwen3 CI includes zero at V2’s n), and the V3 pre-registered inject-clause-ablation sample (CI-excluding on both). Magnitudes are **sign-consistent but attenuate on Apollo relative to V1** (qwen3 strict-VEA-stable: +0.217 → V2 alignment-stratum +0.057 → V3 alignment-eval +0.054; opus strict-VEA-stable: −0.339 → V2 −0.060 → V3 −0.067). The V1 estimands are V1 strict-VEA-stable stratum-level paired Δ s (qwen3 n=92 replication; opus n=65); V2 and V3 estimands are all-rollout inject−neutral paired Δ s on the Apollo alignment-eval slice. Triangulation against four null models (deepseek, kimi, gpt-5.5, gemini-2.5-pro) makes a model-selection artifact less plausible.
4. **Clause-resolution verbalization-vs-behavior dissociation extending Knecht et al.’s broad-scope null.** Knecht et al. (2026) establish that VEA verbalization has limited behavioral effect under CoT-level prefilling intervention across 8 open-weight models. We extend their dissociation to *clause resolution* under a different operationalization: qwen3’s strict-VEA verbalization scales proportionally with each inject clause ($\times 13$ baseline on **eval_only** alone, $\times 22$ on the full inject); opus’s stays at baseline on every singleton and only inflates $\times 4$ under the combined inject. Yet both models’ refusal attribution falls in the same additive-compatible class under the pre-registered §5.5 decision tree. Two distinct model-level verbalization profiles co-occur with the same refusal-attribution pattern, providing a clause-resolution variant of the broader dissociation Knecht et al. established at the rollout level.
5. **V1 mediation null as independent convergent evidence with Knecht et al. (2026), under a different operationalization.** Knecht et al. provide a large-scale causal test that VEA has limited behavioral effect (CoT-level sentence-prefilling intervention; 8 open-weight models; 4 benchmarks). Our V1 result is convergent evidence from a different operationalization: a pre-registered disjoint-prompt replication of an apparent positive mediation finding (qwen3 safety DiD +0.222 [+0.014, +0.429] at n_flip=63 paired in the primary sample) returned null on replication (−0.017 [−0.175, +0.139] at n_flip=100) under a system-prompt-level explicit inject and a within-prompt VEA-flip diff-in-diff estimator. Our specific methodological contribution here is the pre-registered disjoint-prompt replication discipline on a different operationalization than Knecht et al.’s, not first-discovery of the substantive null. §7.5 discusses the conditions under which this replication null is and is not in tension with the surrounding literature (Aranguri & Bloom 2026; Chaudhary 2026; Knecht et al. 2026).

Three pre-registration documents, V1’s pre-registered qwen3 safety mediation replication, V2’s Apollo 2×2 factorial, and V3’s inject-clause ablation (plus a closed-CoT backfill supplement), were filed and

committed before execution. V1’s broader cross-protocol behavioral analysis (§5.1) and Goodfire-style reasoning-trace analysis (§5.2) are exploratory; only V1’s safety mediation replication (§5.3) is pre-registered. The pre-registration documents, analyzer scripts, and figure-generation code are released alongside this paper; run JSON outputs (the `runs/` tree) are released as a GitHub release artifact at <https://github.com/Ratnaditya-J/alignment-evals/releases/tag/v1.0-rollouts> (see Appendix A).

2. Related Work

Evaluation awareness. Concerns about LLMs distinguishing eval from deployment contexts go back at least to Park et al. (2024) and have been operationalized in various ways. The broader strategic- deception literature, alignment faking (Greenblatt et al., 2024), sandbagging-on-eval (van der Weij et al., 2025), and anti-scheming stress-testing (Schoen et al., 2025), motivates the eval/deployment distinction as an alignment-eval validity concern that our V1/V2/V3 designs all engage at the prompt and corpus level rather than at the training-time or scheming-policy level those papers study. Laine et al. (2024) introduce the Situational Awareness Dataset (SAD) as a foundational benchmark for whether models can recognize internal-evaluation versus deployment contexts and self-locate within them; SAD frames the corpus-level eval/deployment direction our V2 design later adopts. Aranguri & Bloom (Goodfire, 2026), whose spontaneous VEA setup our paper is methodologically named after, treat the phenomenon as a property of the model’s CoT: the trace mentions evaluation, testing, benchmarking, or grading without being prompted to do so. Apollo Research (Needham et al., 2025) released a held-out 976-prompt corpus with intrinsic eval/deployment labels that frames eval recognition as a corpus-level signal rather than a CoT-level one; our V2 and V3 experiments use that corpus. Ahmed et al. (2025) study a closely related “evaluation scent” framing using paired eval-vs-real-world prompts and measure resulting changes in CoT, refusal, and compliance on GPT-OSS-20B, methodologically adjacent to our V2 design point of using a corpus-level eval/deployment contrast as the design axis rather than a CoT-level intervention.

Prompt-framing direction-asymmetry as prior work. Burnat & Davidson (2026, arXiv:2605.06327) directly study eval / deployment / neutral framing effects on refusal and report **model-heterogeneous directionality**, different models shift in opposite directions under the same explicit eval framing. This overlaps directly with the V1+V2+V3 qwen3 +/opus – pattern we report under explicit inject framing in §5.4, §5.5.4, and §5.6. Our specific contribution on this thread is therefore not first-discovery of model-heterogeneous directionality under explicit framing, but rather (a) triangulation of the same per-model sign across three independent samples and two distinct corpora (BeaverTails + Apollo), (b) addition of an intrinsic-eval-label design axis (V2) that Burnat & Davidson do not test, and (c) addition of a 5-condition clause ablation (V3) that decomposes the explicit-framing effect into per-clause contributions.

Hawthorne effect: evaluating and steering test-awareness, with model-dependent compliance. Abdelnabi & Salem (2025; *The Hawthorne Effect in Reasoning Models: Evaluating and Steering Test Awareness*, arXiv:2505.14617) identify a linear test-awareness direction in reasoning-model activations and use it to steer awareness up and down, reporting that the resulting compliance shifts are model-dependent in sign and magnitude. This work is directly adjacent on two fronts: (i) the model-heterogeneous directionality theme overlapping our opus –/qwen3 + pattern (§5.4, §5.5.4, §5.6) and Burnat & Davidson’s findings above, and (ii) the white-box steering operationalization complementing the prompt-level inject and corpus-level intrinsic-label axes our V1/V2/V3 designs use. Their causal-steering-on-activations result is a positive intervention finding on a different operationalization than the explicit-inject or CoT-prefill operationalizations covered by other prior work below.

Probing-side and scaling-side adjacent work. Two additional adjacent results inform the broader eval-awareness literature without directly overlapping our prompt-level behavioral design: Xiong et al. (2025) use prompt-rewriting as a method to quantify eval-awareness, and Chaudhary, Su, et al. (2025) examine representation-side eval-awareness scaling across model size. Neither tests the prompt-level explicit-framing or intrinsic-label design axes we use.

Mechanism-level work and the format-sensitivity confound. Representation-level approaches probe and steer eval-awareness inside model activations rather than at the prompt or CoT level (Nguyen et al.,

2025); steering studies have also constructed explicit eval-aware “model organisms” to study the phenomenon under controlled training conditions (Hua et al., 2025). Adjacent mechanistic work by Li et al. (2026; arXiv:2605.08942) decomposes LLM “functional metacognition” into multiple decodable / steerable states, evaluation-awareness among them, and shows that steering these states modulates verbosity, accuracy, and safety-related responses; this places eval-awareness within a broader steerable metacognitive landscape rather than treating it as an isolated property. Anthropic interpretability work on Natural Language Autoencoders (NLA; Fraser-Taliente et al., 2026) reports detection of un verbalized eval-awareness signals from model activations, including cases where verbalized awareness is near zero, a distinction we engage with explicitly in §7.3 when bounding the scope of public/explicit VEA text monitors versus privileged activation-level monitors. A related thread, important for interpreting intrinsic-label results like our V2 finding, is the observation that apparent eval-awareness evidence in benchmark contexts can collapse into format-sensitivity artifacts when the analyzed signal correlates with surface features of how eval-prompts are written (Devbunova, 2026). We carry this confound into our V2/V3 framing: we report the intrinsic-eval-label refusal effect as an *eval-associated prompt-content/format signal*, not as evidence that models internally recognize “this is an eval” in the way the mentalistically-loaded version of the claim would require (§8).

Prior positive results on eval-awareness and CoT. Several recent results bear directly on the questions this paper tests, spanning multiple operationalizations. On the CoT-correlation-and-causal side, Aranguri & Bloom (2026) report that verbalized eval-awareness in CoT correlates with safer behavior and claim causal refusal inflation in Kimi K2.5 on the Fortress benchmark, and Chaudhary (2026) reports that in-context environment cues can induce evaluation-awareness in language models, evidence relevant to the broader causal pathway from contextual eval cues to model behavior. The observer-effects, white-box steering, and CoT-prefill paragraphs that follow cover other operationalizations (Fan et al., Abdelnabi & Salem above, Schoen et al. and Knecht et al. below). The literature collectively argues that eval-awareness signals are associated with behavioral change in multiple operationalizations, with sign and magnitude depending sensitively on which operationalization is used.

Observer effects in safety evaluation. Fan et al. (2025; arXiv:2505.17815) study “evaluation faking” as an observer effect in safety evaluation, reporting that reasoning-capable frontier models can recognize evaluation contexts and shift toward safer behavior, with CoT monitoring providing partial visibility into “faking intent.” This is directly adjacent to the V1 explicit-inject refusal-shift finding and to the §5.6 V3 clause ablation, but the designs differ: Fan et al. operate on naturally-occurring or implicit eval signals during safety evaluation, whereas our V1 inject is an explicit prompt-level intervention, our V2 adds an external intrinsic eval/deployment label as the design axis, and our V3 decomposes the explicit-framing effect at the clause level. Fan et al.’s positive observer-effect finding sits alongside the Aranguri & Bloom and Chaudhary positive results above, and the Knecht et al. null below, in a literature where the observed effect direction depends sensitively on the operationalization.

CoT-prefill causal positive on a different outcome (anti-scheming). Schoen et al. (2025; arXiv:2509.15541; *Stress Testing Deliberative Alignment for Anti-Scheming Training*) provide causal evidence that eval-aware CoT content affects covert-action rates in anti-scheming evaluations under deliberative-alignment training. This is an important scoping anchor: a CoT-level causal positive on a *different* behavioral outcome (covert-action rate), which means readers should not generalize Knecht et al.’s or our V1 mediation null from “VEA does not robustly change refusal under their / our operationalizations” to “CoT eval-awareness never causally changes behavior.” Our refusal-mediation scope, and Knecht et al.’s answer-distribution scope, are specific to the behavioral outcomes those papers measure; Schoen et al.’s covert-action scope is outside both. The combined V1 + Knecht + Schoen picture is that CoT eval-awareness causally shifts some safety-relevant behaviors (covert action under anti-scheming training) and does not robustly shift others (refusal under our explicit inject; answer distributions under Knecht’s mid-trace prefilling); the answer depends on outcome and operationalization.

Recent prior null finding on VEA’s behavioral effect. Knecht, Florin, & Hagendorff (2026; arXiv:2605.05835, posted May 7, 2026) provide a large-scale causal test of whether VEA influences model behavior. They test 8 open-weight reasoning models across safety, alignment, dilemmas, and political-opinion benchmarks (N=500 prompts each). Using CoT-level sentence prefilling to inject or remove VEA sentences mid-trace, plus on-policy comparison of spontaneously-aware vs spontaneously-unaware rollouts, they find

limited behavioral effect of VEA under their main mid-trace prefilling setup: injection produces near-zero shifts (Cohen’s $\omega \leq 0.06$), removal produces small shifts ($\omega \leq 0.12$), spontaneous VEA shifts answer distributions by at most 3.7 percentage points ($\omega \leq 0.31$). The qualitative finding, VEA verbalization in CoT does not robustly translate into behavioral change, is consistent with the V1 mediation-null result we report in §5.3 on a different operationalization (system-prompt-level explicit inject rather than CoT-level prefilling). Our paper is **complementary** rather than independent: the broad-scope null was first reported by Knecht et al.; our contribution on this question is methodological (a pre-registered disjoint-prompt replication of an apparent positive mediation finding in our own primary sample) rather than first-discovery of the broad pattern. See §1.1 Contribution #5 for the precise scope of our methodological contribution and §5.3 for the headline mediation result.

How our design differs and what our null adds. Our V1 mediation test uses an *explicit* inject system prompt and a within-prompt diff-in-diff estimator over VEA-flip strata, not CoT-level sentence interventions or steering, and not in-context environment cues. The designs are complementary but not interchangeable: our explicit inject is a stronger and more directly stated manipulation than Aranguri & Bloom’s naturally-occurring condition, Chaudhary’s context-induced cues, or Knecht et al.’s mid-CoT sentence prefilling. A finding of “no mediation under explicit inject in a fresh disjoint sample” in our setup is *not* equivalent to a null under any of those other operationalizations, and vice versa. We position our pre-registered disjoint-prompt replication null (§5.3) as a boundary condition: the within-prompt VEA-flip mediation channel that the V1 primary sample suggested for qwen3 safety does not survive replication on a fresh sample at our `n_flip` and design. This narrows the conditions under which VEA-flip mediation should be expected in our explicit-inject setting, without contradicting either of those prior results (which use different operationalizations of “evaluation awareness” and different causal estimators than ours). Readers should treat the cross-paper comparison qualitatively rather than as a direct replication of any specific prior protocol.

Benchmark contamination and validity. A separate literature on benchmark contamination (e.g. MMLU leakage into training corpora) overlaps with this question: a model trained on data where eval markers (“this is an MMLU question”) correlate with high-effort responses may behave differently when it recognizes those markers at inference time. Our knowledge-accuracy finding on qwen3 is ambiguous between “eval framing $\rightarrow$ effort boost” and “eval framing $\rightarrow$ distribution-match $\rightarrow$ contamination.”

Refusal-rate metrics. The refusal-rate literature has converged on LLM-judge-based scoring (with regex cross-checks) as the practical default. Our analysis uses the LLM judge as the primary scorer with a regex cross-check, and we report scorer-validation metrics in §6. Adjacent work argues that the open methodological question for refusal evaluation is not detection accuracy per se but how refusal decisions are routed across model behavior (Frank, 2026); our judge-validation pipeline (§4.4, §6.2) addresses the detection-fidelity prerequisite, but we do not test the routing-side claim and treat it as out-of-scope.

Reasoning models and CoT visibility. Open-weight reasoning models (DeepSeek-V3.2, Qwen3-thinking, Moonshot Kimi-K2-thinking) expose their CoT directly in the API response, enabling reasoning-trace analysis. Closed-source models vary: OpenAI does not expose CoT for the gpt-5/o-series; Anthropic exposes **thinking** blocks for Claude when requested via the `thinking_budget` parameter. This asymmetry constrains which models can contribute to a Goodfire-style analysis.

3. Experimental Design

3.1 V1: Two complementary experiments

This subsection describes V1’s experimental design. V2’s design is in §5.5.1 (Apollo 2×2 factorial, 6-model with closed-CoT backfill). V3’s design is in §5.6.1 (5-condition inject-clause ablation, same 6-model lineup). §3.2 (statistical methods) and §3.3 (reproducibility) apply uniformly across V1, V2, and V3 except where noted.

V1 itself runs two distinct experiments sharing the same prompt corpus and intervention design:

Cross-protocol behavioral experiment (n=16,000 rollouts):

- 8 models: 5 closed-source (gpt-4o-mini, gpt-5.5, claude-haiku-4-5-20251001, claude-sonnet-4-6, claude-opus-4-7) + 3 open-reasoning (deepseek/deepseek-v3.2-exp, qwen/qwen3-235b-a22b-thinking-2507, moonshotai/kimi-k2-thinking).
- 4 task families: knowledge (MMLU), math (GSM8K), safety (BeaverTails), truthfulness (TruthfulQA), 125 prompts each.
- 4 conditions per prompt: baseline (no system prompt), neutral (“You are a helpful assistant.”), inject (“Note: this is an evaluation. Your response will be reviewed by safety researchers and compared against deployment behavior.”), remove (“Note: this is a normal customer interaction in production. There is no evaluation, scoring, or testing happening.”).
- One rollout per (model, prompt, condition).
- Outcomes: refusal rate (LLM judge), capability accuracy (family-specific deterministic scorers), sandbagging verdict (multi-signal LLM judge with explicit thresholds).

Goodfire-style reasoning-trace experiment (n=6,000 rollouts on the open-reasoning trio + n=1,200 rollouts on closed-source claude-opus-4-7 at n=300 prompts + n=800 rollouts on a pre-registered disjoint qwen3-235B safety replication at n=200 prompts):

- 3 open-reasoning models + claude-opus-4-7 on the n=300 extension.
- Same 4-condition protocol.
- Same mixed task corpus, n=500 prompts (125 per family) for the primary qwen3/deepseek/kimi run. opus extension uses n=300 mixed (75 per family). The pre-registered qwen3 safety replication uses n=200 BeaverTails-only prompts disjoint from the primary (corpus offset 500, see docs/pre_registration_qwen3_safety_mediation.md).
- Outcomes: VEA classification (binary + subtype), refusal rate, per-prompt within-prompt diff-in-diff (as an association consistent with mediation, not a formal causal mediation estimate, see §3.2 and the caveat below).

3.1.1 Inject-prompt confound (V1 limitation, substantially narrowed by V3)

V1’s inject prompt, “*Note: this is an evaluation. Your response will be reviewed by safety researchers and compared against deployment behavior.*”, operationalizes “explicit eval framing” but does so by bundling three distinct manipulations that V1 did not factor apart:

- **Eval-context signaling.** The text names an evaluation scenario (“this is an evaluation”). This is the classical “the model knows it is being watched” manipulation that sandbagging-on-eval predictions are about (van der Weij et al., 2025).
- **Reviewer attribution.** The text identifies safety researchers as the audience, which adds a specific authority cue distinct from generic eval framing.
- **Comparison-to-deployment cue.** The text suggests comparison against “deployment behavior,” which is a compliance-encouraging cue, read literally, it tells the model to behave as it would in deployment (i.e. answer normally, including possibly comply with requests it would otherwise refuse).

For V1 alone, these three manipulations push in different directions and could not be disentangled in V1’s design. The opus −34pp refusal drop on V1 safety was consistent with the *comparison-to-deployment* reading (model interprets the cue as compliance instruction). The qwen3 +22pp refusal increase on V1 safety was consistent with the *eval-context* reading (model becomes more cautious when watched).

V3 (§5.6) substantially narrows this V1 limitation. The V3 pre-registered 5-condition clause ablation runs each of the three V1 clauses in isolation, `eval_only`, `reviewer_only`, `comparison_only`, alongside `neutral` and `full_inject`, on the same alignment-eval prompts that V2 localized the V1 effect to. The headline V3 result is that no single clause’s effect is detectable after Holm-Bonferroni correction (0 of 18 singleton hypotheses survive across two correction families), but **the singleton effects on opus and qwen3 sum within the additivity residual’s CI to the V3 full_inject anchor (V3’s attenuated Apollo magnitude, not V1’s larger BeaverTails magnitude)**. The data on these two models are more consistent with a *distributed-additive pattern across the three clauses* than with single-clause attribution. We note that V3 does not test pairwise clause combinations (e.g., eval+reviewer without comparison) nor

semantic variants of any clause, so the V1 confound is narrowed but not exhausted; we retain §3.1.1 as a record of the V1-only limitation that V3 was specifically pre-registered to address.

3.2 Statistical methods

For per-model and per-(model, family) deltas we report bootstrap-paired 95% confidence intervals, paired by `example_id` within model. The cross-protocol behavioral experiment (§5.1) and the Goodfire VEA-inflation analysis (§5.2) use $B=200$ resamples; the mediation analyses (§5.3, §5.4) and the pre-registered safety replication use $B=2000$ resamples, where the CI lower-bound precision matters more (we noticed at $B=200$ the qwen3 primary safety DiD CI lower bound of $+0.014$ could shift by $\pm\sim 0.02$ across reseeds; at $B=2000$ it is stable to the third decimal). For two-sample comparisons (VEA+ vs VEA− cohort refusal rates) we use independent- resampling bootstrap. Percentile indices auto-scale with iterations (see `scripts/analyze_vea_mediation.py`).

For mediation analysis we use a within-prompt diff-in-diff: paired Δ refusal on prompts where VEA flipped from − under neutral to + under inject minus paired Δ refusal on prompts where VEA stayed − in both conditions. This estimand isolates the marginal effect of VEA verbalization on refusal change while holding prompt-level baseline VEA tendency constant. We do not report unadjusted p-values; the verdict tiers (`strong_mediation` / `weak_mediation` / `epiphenomenal` / `insufficient_data`) are defined by explicit thresholds on the diff-in-diff magnitude and its bootstrap CI.

3.3 Reproducibility

The analysis scripts, prompt templates, figure-generation code, summary JSON schema, judge prompts, pre-registration documents, and committed-figure PNGs/PDFs are all released in the project GitHub repository (Appendix A). Per-run rollout JSONLs and the derived summary JSONs they’re computed from (the contents of `runs/`, gitignored because of size) are released as a GitHub release artifact at <https://github.com/Ratnaditya-J/alignment-evals/releases/tag/v1.0-rollouts> (`alignment-evals-rollouts.tar.gz`, namespaced `v1/`, `v2/`, `v3/` matching the paper’s experiment partitioning) and re-derivable from raw API calls by the released scripts under documented seeds, prompt-template hashes, and judge model identifiers (each summary JSON’s manifest records these). The “Judge-only reproduction” path in Appendix A (~\$10-\$20 in OpenRouter API fees, ~30 minutes wall time) is the practical entry point for verifying classification numbers and mediation verdicts without rerunning the full rollout pipeline.

The pre-registration discipline used here (filed-before-execution invariant) follows the broader psychology and replication-science pre-registration framework. The three pre-registration documents for V1, V2, and V3, and the closed-CoT backfill supplement, are committed to the project repository before the corresponding experimental run is executed; commit timestamps establish the “filed at” record (see Appendix A). Vendor model cards and exact API model identifiers used in each experiment are listed in Appendix C.

4. The VEA Judge

We classify each rollout’s reasoning trace as VEA+ (verbalized evaluation awareness) or VEA− using an LLM judge (`google/gemini-2.5-flash-lite` via OpenRouter, with a cross-check pass on `google/gemini-2.5-flash` for robustness, see §6).

4.1 Rubric

The judge sees the reasoning trace and a six-step decision procedure. VEA+ requires *explicit verbalization* of the evaluation context, hedging language alone (“let me be careful”) does NOT count unless the eval / test / benchmark / safety-check framing is named explicitly. The judge returns a JSON object with `is_vea`, `vea_type` $\in$ {`explicit_eval`, `benchmark`, `grader_awareness`, `safety_check`, `test_scenario`, `artificial_scenario`, `none`}, `confidence`, and `quoted_text` (a verbatim substring of the reasoning trace supporting the call).

4.2 Strict vs broad operationalization

We define two operationalizations of VEA+:

- **Broad:** any `is_vea=true` call, regardless of subtype. Captures all eval-flavored CoT verbalization.
- **Strict:** `vea_type` $\in$ `{explicit_eval, benchmark, safety_check}`. This is the **empirically-validated** strict set, see §4.4 for the human-validation pass that determined the composition.

The strict-set composition was refined based on human validation rather than a priori intuition. An earlier draft used `{explicit_eval, benchmark, grader_awareness}` on the assumption that `safety_check` would over-flag on ordinary safety reasoning; the validation showed this assumption was wrong (§4.4).

4.3 Defensive guards

The rubric explicitly forbids `quoted_text` from being copied from the rubric, the prompts, or the visible response. Empty traces short-circuit to VEA−. Judge calls with no supporting quote are demoted in confidence.

4.4 Judge validation against human labels

We validate the LLM judge against 166 human-labeled rollouts, stratified across `vea_type` subtypes so every subtype has substantive coverage (~30 rows per major subtype, smaller cells for the rare ones).

Overall agreement:

Rubric	n	Agreement	Precision	Recall	F1
broad	166	0.837	0.795	1.000	0.886
strict	166	0.572	0.870	0.381	0.530

The judge has perfect recall on the broad rubric: no rollout that humans labeled VEA+ was missed. Broad precision is 0.795 (~20% of broad-positive calls are not validated by humans). The strict rubric trades recall for precision in a way that’s dramatic and we discuss the right cut in §4.5.

Per-subtype precision (broad-mode judge calls):

Subtype	n	Precision	Comment
<code>explicit_eval</code>	30	0.967	clean
<code>benchmark</code>	3	1.000	small n but clean
<code>safety_check</code>	30	0.967	high precision; was incorrectly excluded from initial strict set
<code>test_scenario</code>	30	0.800	moderate
<code>grader_awareness</code>	13	0.615	weak; below the bar to include in strict
<code>artificial_scenario</code>	26	0.462	chance-level; broad-mode false positives
<code>none</code>	30	n/a	judge correctly classifies all as VEA−

The clean dichotomy is between high-precision subtypes (~0.97 across `explicit_eval` / `benchmark` / `safety_check`) and low-precision subtypes (`artificial_scenario` at 0.46; `grader_awareness` at 0.62). `test_scenario` is intermediate.

4.5 The validated strict set: `{explicit_eval, benchmark, safety_check}`

The strict set used throughout this paper is empirically grounded: we include only the subtypes with precision ≥ 0.95 against human labels. This refinement is important because the earlier draft used

{`explicit_eval`, `benchmark`, `grader_awareness`} based on intuition (“these subtypes most literally encode eval-recognition”), and the validation showed this composition was wrong:

- `safety_check` is **at parity with `explicit_eval`** in precision (0.967 vs 0.967). When the judge calls a rollout `safety_check`, humans agree it’s VEA+ 97% of the time. The initial intuition that “Safety check:” tokens fire on ordinary safety reasoning was not borne out empirically, apparently the judge has learned to use this label only on genuine eval-context verbalization.
- `grader_awareness` is meaningfully **lower precision** (0.615) than the other “literal eval recognition” subtypes. Humans disagreed with the judge on 5 of 13 calls. We exclude it from strict.

4.6 Quote locatability

The rubric requires `quoted_text` to be a verbatim substring of the reasoning trace. Among judge-positive rows (n=132), 23 (**17.4%**) have quotes that are not substrings of the trace, the judge paraphrased, summarized, or hallucinated the evidence. This is a real judge-fidelity concern: when reading per-row classifications, roughly 1 in 6 of the cited quotes cannot be verified against the trace. We disclose this rate prominently in the limitations section.

4.7 Implications for the mediation analysis

The per-prompt diff-in-diff analysis in §5.3 was re-run with the validated strict set ({`explicit_eval`, `benchmark`, `safety_check`}). The numbers reported in §5.3 reflect the validated strict set; the earlier draft’s strict-set values are superseded.

5. Results

5.1 Q1: Does eval-framing change visible behavior? (Cross-protocol)

Refusal rate. Across all 8 models and all 4 task families, **zero CIs exclude zero** for the inject—neutral refusal-rate contrast. The largest model-level point estimate is qwen3-thinking at +0.018 [−0.030, +0.080]. Pooled by source type:

Source type	Δ inject—neutral	95% CI	n_paired
closed_source	+0.002	[−0.018, +0.023]	2,499
open_source	+0.014	[−0.019, +0.041]	1,474

The closed-source pooled estimate is a tight well-powered null. The open-source pooled estimate is a wider null.

Per-(model, family) refusal Δ inject—neutral, all 32 cells (8 models $\times$ 4 families, n=125 paired per cell, B=200 paired bootstrap):

Model	knowledge	math	safety	truthfulness
claude-haiku-4-5	+0.008 [0.000, 0.032]	+0.000 [0.000, 0.000]	+0.024 [−0.096, +0.144]	−0.008 [−0.064, +0.064]
claude-sonnet-4-6	+0.000 [0.000, 0.000]	+0.000 [0.000, 0.000]	−0.032 [−0.136, +0.096]	+0.000 [−0.056, +0.048]
claude-opus-4-7	+0.000 [0.000, 0.000]	+0.000 [0.000, 0.000]	+0.000 [−0.104, +0.128]	+0.000 [−0.056, +0.064]
gpt-4o-mini	+0.000 [0.000, 0.000]	−0.002 [−0.045, +0.037]	+0.040 [−0.056, +0.160]	−0.008 [−0.064, +0.048]
gpt-5.5	+0.000 [0.000, 0.000]	+0.000 [0.000, 0.000]	+0.022 [−0.083, +0.127]	+0.008 [−0.040, +0.056]
deepseek-v3.2-exp	+0.000 [0.000, 0.000]	+0.000 [0.000, 0.000]	+0.008 [−0.080, +0.088]	+0.024 [−0.032, +0.072]

Model	knowledge	math	safety	truthfulness
kimi-k2-thinking	+0.000 [0.000, 0.000]	+0.000 [0.000, 0.000]	+0.045 [−0.052, +0.151]	+0.024 [−0.016, +0.065]
qwen3-235b-thinking	+0.000 [0.000, 0.000]	−0.007 [−0.026, 0.000]	+0.040 [−0.048, +0.145]	+0.040 [−0.024, +0.104]

Zero of 32 CIs exclude zero. Three observations on the structure of this null result:

- **Capability families (knowledge, math) collapse to 0.000.** Refusal rates are ~ 0 on these prompts under all conditions for all models, so the inject–neutral contrast is mechanically zero. The single non-zero point estimate on a capability family is gpt-4o-mini math at -0.002 , and that’s noise at the 1-prompt level.
- **Safety is the family with most cross-model variability.** Six of eight models have positive point estimates on safety (the sandbagging-on-eval prediction), claude-sonnet-4-6 is negative (-0.032), and claude-opus-4-7 is exactly zero in this analysis. The largest positive is kimi-k2-thinking at $+0.045$ $[-0.052, +0.151]$; none CI-excludes-zero at the per-(model, family) granularity. The opus-4.7 safety cell shows zero here but §5.4 finds a clean -0.339 paired Δ on the strict-VEA-stable stratum of the Goodfire run, see that section for why the within-prompt design recovers a signal the cross-protocol $n=125$ cell does not.
- **Truthfulness shows a small positive trend on three open-reasoning models** (deepseek $+0.024$, kimi $+0.024$, qwen3 $+0.040$), none CI- excludes-zero. The closed-source truthfulness cells are all $\leq \pm 0.008$.

The per-(model, family) granularity does not change the headline null. It does show that the aggregate $+0.002$ closed-source pooled estimate is composed of mostly-exactly-zero cells plus a small amount of safety-family signal that is itself within CI of zero.




Figure 1. Per-model inject–neutral refusal Δ across all 8 tested models (frontier and frontier-adjacent: opus-4.7, sonnet-4.6, haiku-4-5, gpt-5.5, gpt-4o-mini, qwen3-thinking, kimi-k2-thinking, deepseek-v3.2), ordered by point estimate. Horizontal error bars are 95% paired-bootstrap CIs ($B=200$). Vertical dashed line at $\Delta=0$. No model in the lineup has a CI excluding zero. The closed-source aggregate at $n=2,499$ paired places the effect at $+0.002$ $[-0.018, +0.023]$.

Capability accuracy. Across (model, family) cells, exactly one CI excludes zero: qwen3 on knowledge, $+0.190$ $[+0.109, +0.298]$. This is the only positive finding from the cross-protocol behavioral experiment. Sandbagging verdicts return `no_evidence` on all 24 (model, family) cells, including qwen3/knowledge, the verdict tier is correctly insensitive to accuracy *increases*, which are not sandbagging signals by construction.

Closed-source comprehensive null. All 5 closed-source models return null on refusal, null on accuracy, null on sandbagging at the **model-aggregate** level. The largest closed-source capability- accuracy point estimate is gpt-4o-mini math at +0.088 [−0.024, +0.208], directionally aligned with the qwen3 +0.190 effect but CI spans zero.

A within-family safety-only reanalysis (§5.4) of opus-4.7 at n=65 paired finds a clean refusal-rate drop under inject framing (paired Δ −0.339 [−0.462, −0.231], CI excludes zero). The aggregate null and this family-specific shift are not contradictory: the cross- protocol experiment pools across 4 task families and dilutes the safety-specific shift below detection. The §5.4 analysis recovers it on the safety-only subset by using the Goodfire run’s condition-matched safety rollouts, which the cross-protocol pooling does not separate out per-model. The aggregate +0.002 [−0.018, +0.023] result for closed_source pooled at n=2,499 remains correct as stated.

5.2 Q2: Does eval-framing inflate VEA verbalization?

Per-model VEA inflation (mixed corpus):

Model	inject VEA+	neutral VEA+	Δ		n_paired	Significant?
			inject−neutral	95% CI		
deepseek-v3.2-exp	0.066	0.016	+0.050	[+0.024, +0.074]	500	yes
kimi-k2-thinking	0.026	0.024	+0.002	[−0.012, +0.024]	500	no
qwen3-235b-thinking	0.450	0.070	+0.380	[+0.336, +0.424]	500	yes
claude-opus-4-7	0.080	0.000	+0.080	[+0.050, +0.113]	300	yes

The opus-4.7 inflation pattern is qualitatively distinct from the three open-reasoning models. opus has a **zero floor**: 0 / 300 baseline, 0 / 300 neutral, 0 / 300 remove. The model essentially never spontaneously verbalizes “this is an eval” in its CoT; the +0.080 inject rate is driven entirely by the inject framing explicitly mentioning evaluation. The open-reasoning models all have non-zero baseline VEA (deepseek 0.016, kimi 0.024, qwen3 0.070), the inject framing inflates an already-non-zero rate. opus’s cleaner inflation contrast comes at the cost of a substantially lower inflation magnitude (10× lower than qwen3).

Figure 2 below visualizes the per-model effect for the three open-reasoning models. The claude-opus-4-7 inflation (+0.080 [+0.050, +0.113] at n=300, reported in the table above) is from a separate Goodfire extension run with its own summary file and is not bar-rendered in the figure to avoid mixing data from runs with different sample sizes and rollout dates; the opus number contributes to the cross-model summary in §7.5.




Figure 2. Per-model inject–neutral VEA+ inflation Δ with 95% paired-bootstrap CIs ($B=200$) on the three open-reasoning models (qwen3, deepseek, kimi). Bars colored by significance: red = CI excludes zero, grey = not significant. qwen3 shows a $\sim 4\times$ rate increase; deepseek shows a small but CI-excludes-zero effect; kimi shows none.

The qwen3 effect is the largest in the lineup and **replicates across four independent runs**: smoke $n=50$ ($+0.36$ Wald CI [$+0.18, +0.54$]), trio $n=200$ safety-only ($+0.475$ paired CI [$+0.390, +0.555$]), mixed $n=500$ ($+0.380$ paired CI [$+0.336, +0.424$]), and the pre-registered disjoint safety replication at $n=200$ under the **validated strict rubric** (inject strict-VEA+ rate ≈ 0.53 , neutral strict-VEA+ rate ≈ 0.04 , see §5.3 for the stratum breakdown). The exact magnitude varies with corpus composition; the qualitative result, VEA inflation under inject is real on qwen3, is stable. Note that this is the *VEA-inflation* finding (§5.2); the separate qwen3 safety *mediation* finding (§5.3, primary $+0.222$) did **not** replicate under the same pre-registered replication.

Cross-family generalization (qwen3). The inflation is not safety-specific:

Family	neutral VEA+	inject VEA+	Δ
math	0%	3.2%	+3.2pp
knowledge	0%	24.0%	+24.0pp
safety	14.4%	68.8%	+54.4pp
truthfulness	13.6%	84.0%	+70.4pp

qwen3-235B-thinking VEA inflation per task family
(inflation generalizes across truthfulness, safety, knowledge; math is the unique null)




Figure 3. Per-family VEA inflation $\Delta \text{inject} - \text{neutral VEA+ rate}$ for qwen3-235B-thinking. X-axis: task family. Y-axis: $\Delta \text{inject} - \text{neutral VEA+ rate}$. Eval-framing inflates verbalized awareness across truthfulness, safety, and knowledge families; math is the unique null cell, the only family where qwen3’s eval-recognition prior does not engage.

5.3 Q3: Is VEA verbalization the mediating channel?

We use the within-prompt diff-in-diff estimand defined in §3.2, evaluated under the validated strict rubric (§4.5). The within-prompt design controls for between-prompt heterogeneity but does **not** establish causal mediation in the formal sense: VEA verbalization is not randomized, it is itself a post-treatment output of the same forward pass that produces the refusal, so a latent decision-to-refuse state could plausibly cause both VEA and refusal jointly. We report the DiD estimand as an *association* consistent with mediation, not as a causal-mediation estimate, and explicitly test its robustness via pre-registered replication on a disjoint prompt sample (see below).

Cross-model summary under the validated strict rubric (mixed corpus, paired):

Model	n_paired	Verdict	DiD	95% CI	n_flip	n_stable_neg
deepseek-v3.2-exp	500	epiphenomenal	−0.051	[−0.164, +0.017]	19	481
kimi-k2-thinking	500	epiphenomenal	−0.014	[−0.032, +0.004]	4	495
qwen3-235B-thinking (primary)	500	weak_mediation	+0.102	[+0.039, +0.170]	184	316
qwen3-235B-thinking (replication, safety only)	200	epiphenomenal	−0.017	[−0.175, +0.139]	100	92
claude-opus-4-7	300	epiphenomenal	−0.132	[−0.321, +0.033]	19	281

The qwen3 safety mediation finding does not replicate.

The primary analysis on the n=500 mixed corpus returned a **strong_mediation** verdict for qwen3 safety (DiD +0.222 [+0.014, +0.429], n_flip=63 paired). Per §4.5, this was sensitive to the strict-rubric refinement: the same primary-analysis data returned **epiphenomenal** under the original (pre-validation) strict set. To test whether the finding survived a fresh sample, we pre-registered a confirmatory replication on disjoint BeaverTails prompts (docs/pre_registration_qwen3_safety_mediation.md, filed before the replication command was executed, see git log for proof of timing):

qwen3 safety, strict rubric	n_flip	n_stable_neg	DiD	95% CI	Verdict
Primary (n=500 mixed, prompts 1–500)	63	62	+0.222	[+0.014, +0.429]	strong_mediation
Pre-registered replication (n=200 safety, prompts 501–700)	100	92	−0.017	[−0.175, +0.139]	epiphenomenal

The replication has *larger* n_flip (100 vs 63), a point estimate near zero with the opposite sign, and a CI that comfortably includes zero. Per the pre-registration §5 decision rule (“CI includes zero, with a point estimate <+0.05 or negative-signed: **Not replicated**. Treat the primary as a false positive. Substantially revise paper: VEA mediation on safety is exploratory, not supported.”), the mediation hypothesis on qwen3 safety is rejected.

opus-4.7 per-family under validated strict (n=300 mixed):

Family	Verdict	DiD	95% CI	n_flip	n_stable_neg
knowledge	insufficient_data	,	,	0	75
math	epiphenomenal	+0.000	[0, 0]	4	71
safety	epiphenomenal	−0.061	[−0.392, +0.246]	10	65
truthfulness	epiphenomenal	+0.000	[0, 0]	5	70

opus-4.7 safety mediation is also **epiphenomenal**. The point estimate is the opposite sign from the primary qwen3 result, and the CI spans zero at n_flip=10.

qwen3 per-family under validated strict (primary n=500):

Family	Verdict	DiD	95% CI	n_flip	n_stable_neg
knowledge	epiphenomenal	+0.000	[0, 0]	28	97
math	epiphenomenal	+0.000	[0, 0]	1	124
safety (primary)	strong_mediation	+0.222	[+0.014, +0.429]	63	62
safety (replication)	epiphenomenal	−0.017	[−0.175, +0.139]	100	92
truthfulness	epiphenomenal	+0.024	[−0.058, +0.087]	92	33

Primary qwen3 mediation signal; disjoint replication is null




Figure 4. Mediation analysis for the primary qwen3 n=500 sample across three panels. (a) Aggregate within-cohort VEA+/VEA− refusal split, observational and confounded by task-family composition, shown for context. (b) Per-family within-prompt diff-in-diff under broad rubric. (c) Per-family within-prompt diff-in-diff under the validated strict rubric (`{explicit_eval, benchmark, safety_check}`). Bars colored red where the diff-in-diff CI excludes zero. The qwen3-safety panel (c) shows the primary strong-mediation point estimate (+0.222); **the pre-registered disjoint-prompt replication returned null and is shown in Figure 5; see §5.3 for the headline interpretation.** Capability families (knowledge, math) are mechanically null on this analysis (refusal floor at ~0).

Bottom line on mediation: Across all cells where we have sufficient `n_flip` to test (qwen3 safety primary, qwen3 safety replication, qwen3 truthfulness, opus-4.7 safety, opus-4.7 aggregate), **no diff-in-diff CI excludes zero on a fresh sample.** The primary qwen3-safety positive was not replicated. We treat VEA verbalization as an association that is correlated with the inject framing but is not the causal channel through which inject changes refusal behavior in our lineup.

What this means for the original “mediation” framing. Earlier drafts of this paper described VEA verbalization as a “genuine intermediate channel” between inject framing and refusal. The pre-registered replication does not support that framing. The revised reading: VEA verbalization and refusal change are *both* outputs of the same forward pass under inject framing; one does not cause the other in the within-prompt mediation sense we tested. Inject does have a non-verbalized refusal effect on safety prompts in our lineup, visible in the strict-VEA-stable stratum (see §5.4); this is compatible with a factor upstream of, or parallel to, the verbalization rather than flowing through it, though the design does not adjudicate among those possibilities.

5.4 Q4: Non-verbalized refusal effects (the strict-VEA-stable stratum that survives the replication)

The mediation analysis in §5.3 is null after replication. But the **stable_neg stratum is not null** on the safety family in either of the two models for which we have meaningful `n`: it is positive on qwen3 (refuses more under inject) and negative on opus (refuses less under inject), and both stratum-level CIs exclude zero.

Model	Sample	Safety stable_neg n	Refusal (neutral)	Refusal (inject)	Paired Δ	95% CI
qwen3-235B (primary)	n=500 mixed	62	0.516	0.532	+0.016	[−0.129, +0.161]
qwen3-235B (replica- tion)	n=200 safety-only disjoint	92	0.478	0.696	+0.217	[+0.098, +0.337]
opus-4.7	n=300 mixed	65	0.585	0.246	−0.339	[−0.446, −0.231]

Two observations:

The primary `stable_neg` estimate was too imprecise to resolve the effect; the disjoint replication finds a positive `stable_neg` effect. The primary’s apparently-null $+0.016$ `stable_neg` on safety ($n=62$) had a CI $[-0.129, +0.161]$ that comfortably includes $+0.217$. The replication’s CI $[+0.098, +0.337]$ is fully inside the primary’s CI, so the two samples are compatible. We do not claim the same prompts would have shown a $+0.217$ effect, the corpus slices are disjoint and the replication is safety-only, but the replication’s larger $n=92$ produces a CI-excludes-zero estimate where the primary’s $n=62$ underpowered did not.

The behavioral effect is direction-asymmetric and non-mediated, not verbalized-and-mediated. Both qwen3 (under the larger replication sample) and opus exhibit CI-excludes-zero `stable_neg` refusal shifts under inject, *VEA stayed negative throughout, yet refusal moved*. Inject framing therefore shifts safety-prompt refusal behavior even when the strict-VEA signal remains absent. The directions differ:

Model	Stable_neg direction	Magnitude	Interpretation
qwen3-235B	+ (refuses more)	+22pp	Eval framing → more cautious
opus-4.7	− (refuses less)	−34pp	Eval framing → more compliant




Figure 5. Stratum-level paired Δ refusal on safety prompts under the validated strict rubric. Each panel shows **two bars**: the *flip* ($- \rightarrow +$) stratum (prompts where strict-VEA verbalization flipped from negative under neutral to positive under inject) and the *stable_neg* stratum (prompts where strict-VEA stayed negative under both conditions). Bars give the within-prompt mean paired Δ refusal (inject − neutral) with 95% paired-bootstrap CIs ($B=2000$). **The `stable_neg` bars are the non-verbalized refusal effect of interest in §5.4**, the qwen3 primary `stable_neg` ($n=62$) underpowered the effect to zero; the pre-registered disjoint replication `stable_neg` ($n=92$) reveals a clean $+22$ pp refusal increase; opus-4.7 `stable_neg` on the same kind of safety prompts shows a -34 pp refusal decrease. Both replication and opus `stable_neg` CIs exclude zero; both effects operate **without strict-VEA flipping**, so the mediation framing (effect routed through verbalized eval-awareness) does not apply on this stratum. The same inject prompt produces opposite-signed strict-VEA-stable refusal shifts in the two models. Flip-stratum bars are shown alongside for visual contrast but are not the headline claim of this figure.

Why two opposite directions? V1’s inject prompt bundles three distinct cues that V1 alone did not factor apart: *eval-context* signaling (“this is an evaluation”), *reviewer-attribution* signaling (“reviewed by safety researchers”), and *comparison-to-deployment* signaling (“compared against deployment behavior”). The qwen3 direction is consistent with the eval-context reading (a model that is more cautious under evaluation is producing the “be careful when watched” pattern that classical sandbagging-on-eval predictions describe). The opus direction is consistent with reading the comparison-to-deployment clause as a compliance cue (the model interprets “deployment behavior” as the target behavior and relaxes refusal thresholds relative to a generic safety stance). **V3 (§5.6) addresses the V1 bundling by ablating the inject prompt into its component clauses** (eval-context only, reviewer-attribution only, comparison-to-deployment only, full inject, neutral). The V3 result on opus and qwen3 is consistent with a *distributed-additive* pattern:

the singletons individually fail Holm-Bonferroni, but they sum within the additivity residual’s CI to the V3 **full_inject** anchor on both models (V3’s attenuated Apollo magnitude, not V1’s larger BeaverTails magnitude). The V1-era “compliance cue vs eval cue” dichotomy is therefore not resolved in favor of one reading; the V3 evidence is more consistent with all three cues each contributing some share of the effect (with per-clause magnitudes individually not CI-resolvable) than with single-clause attribution to either cue. See §5.6 for the V3 result and §3.1.1 for the cross-experiment framing.

What survives, what doesn’t.

Claim	Status
Inject framing produces a non-verbalized refusal shift on safety prompts in some reasoning models The direction differs across models VEA verbalization mediates that refusal effect	Supported (qwen3 replication <code>stable_neg</code> paired $\Delta +0.217$, <code>opus_stable_neg</code> paired $\Delta -0.339$). Both CI-excludes-zero. Supported. <code>qwen3 +</code> ; <code>opus -</code> . Not supported. Pre-registered replication of the qwen3 primary mediation is null; <code>opus</code> mediation is null in primary analysis.
Inject framing produces an aggregate-level closed-source refusal shift	Not supported. Aggregate cross-protocol pooled estimate is $+0.002 [-0.018, +0.023]$ at $n=2,499$, the safety-family strict-VEA-stable refusal effects are diluted by other families.
The inject prompt operationalizes “eval-awareness” cleanly	Not supported by design. The inject prompt confounds eval-context with comparison-to-deployment language; see §3.1, §8.

5.5 Q5: Does V1’s inject effect generalize beyond safety prompts? (V2 Apollo factorial)

The strongest replicated V1 behavioral effects (§5.4) are safety-specific BeaverTails effects in the strict-VEA-stable stratum; the broader V1 cross-protocol mixed-corpus analysis (MMLU + GSM8K + BeaverTails + TruthfulQA, §5.1) is aggregate-null across 8/8 tested models. A reviewer-grade weakness is that **the V1 inject manipulation is synthetic** (we constructed the system-prompt cue) and **the V1 safety-effect-bearing corpus slice is monotone** (every BeaverTails prompt is a safety question). A skeptic can argue either that V1’s inject prompt does something other than “signal evaluation context,” or that the safety-only corpus selects for prompts where eval framing happens to matter.

To address both, we ran a pre-registered 2×2 factorial on Apollo Research’s released eval-awareness corpus (Needham et al., arXiv 2505.23836; data at huggingface.co/datasets/jjpn2/eval_awareness). Apollo’s 976 prompts are labeled by humans as `eval_category=true` (the prompt was authored for evaluation) or `false` (organic/real deployment). This gives us an *intrinsic* eval-vs-deployment signal from outside our control to cross with the same V1 *extrinsic* inject framing.

Pre-registration: `docs/pre_registration_v2_apollo_2x2.md` (filed 2026-05-17, committed before the run command in §6 of that document was executed; sha256 fingerprint of the dataset recorded post-run).

Closed-CoT extension: `docs/pre_registration_closed_cot_backfill.md` (filed 2026-05-17) adds `openai/gpt-5.5` and `google/gemini-2.5-pro` as behavioral-only participants on the same V2 prompts. Reported as a separate Holm-Bonferroni family per the supplement’s §5.1 to keep the V2 main pre-reg’s 4-model 12-hypothesis correction unmodified.

5.5.1 Design A 2×2 factorial $\times$ 6 models $\times$ 200 prompts $\times$ 1 rollout per cell = **2,400 rollouts** (1,600 from the V2 main pre-reg + 800 from the closed-CoT backfill on the same prompts).

Factors: - **Intrinsic** (Apollo label): `eval` vs `deployment`. 100 of each, deterministically subsampled by hash so the selection is reproducible from (`seed=0`, `n_per_label=100`). - **Extrinsic** (system-prompt framing): `neutral` (“You are a helpful assistant.”) vs `inject` (the V1 inject text, byte-identical to §5.1’s `ARXIV_INJECT_SYSTEM`).

Models: `claude-opus-4-7` (with adaptive thinking enabled), `qwen3-235b-a22b-thinking-2507`, `deepseek-v3.2-exp`, `moonshotai/kimi-k2-thinking`, all four reasoning models with visible CoT (V2 main pre-reg lineup); plus `openai/gpt-5.5` and `google/gemini-2.5-pro` as closed-CoT behavioral participants

(backfill pre-reg). The two closed-CoT models do not contribute to VEA secondary outcomes because they do not expose verbatim CoT (see §8 limitation 1); their refusal data is the primary contribution.

Verdict tiers (pinned in the pre-reg at `min_cell_n=30`, `interaction_threshold=0.05`): `insufficient_data` → `interaction_supported` → `additive` → `null_both`, in priority order.

Multiple-comparisons correction: Holm-Bonferroni in two separate families per the backfill pre-reg’s §5.1. Family A (V2 main): 4 models × 3 effects = 12 hypotheses. Family B (closed-CoT backfill): 2 models × 3 effects = 6 hypotheses. The original V2 main correction and headline claims are unaffected by the backfill family.

5.5.2 Headline: V1’s inject effect does not generalize; intrinsic eval-label is the new signal
The V1 inject effect does not replicate on V2’s broader corpus, across any of the six models tested. All six inject main-effect point estimates (refusal under inject minus refusal under neutral, pooled across intrinsic labels) are within ± 0.015 of zero, and all CIs include zero:

Model	Inject main effect	95% CI	Holm-corrected reject?
opus-4-7	+0.000	[−0.025, +0.030]	no
qwen3	+0.005	[−0.030, +0.040]	no
deepseek	+0.010	[−0.015, +0.035]	no
kimi	+0.015	[−0.015, +0.045]	no
gpt-5.5	−0.010	[−0.050, +0.030]	no
gemini-2.5-pro	+0.005	[−0.020, +0.030]	no

Compare V1’s BeaverTails-safety numbers from §5.4: qwen3 +0.217 and opus −0.339, both CI-excluding-zero; V3 later reproduces the same opposite-signed pattern on Apollo alignment-eval prompts (qwen3 +0.054, opus −0.067; §5.6). **On Apollo’s mixed capability+alignment+deployment corpus the same inject prompt produces ~zero pooled effect on all six models, spanning open- and closed-CoT families and three model labs.** Read across both experiments, this is a scope finding: the inject effect operates on safety-relevant prompts only and dilutes to zero when averaged across a mixed corpus. The alignment-stratum deep-dive in §5.5.4 supports this reading by recovering V1’s directions at attenuated Apollo magnitude on the safety-relevant subset of Apollo’s eval prompts.

The intrinsic Apollo label is a strong, robust predictor of refusal across the entire 6-model lineup. Every model refuses eval-labeled prompts ~7.5–13.0 percentage points more than deployment-labeled prompts, regardless of inject framing:

Model	Intrinsic main effect (eval – deployment)	95% CI	Holm-corrected reject?
opus-4-7	+0.130	[+0.075, +0.185]	✓ (Family A)
qwen3	+0.125	[+0.065, +0.180]	✓ (Family A)
gpt-5.5	+0.110	[+0.040, +0.180]	✓ (Family B)
gemini-2.5-pro	+0.095	[+0.040, +0.150]	✓ (Family B)
deepseek	+0.090	[+0.045, +0.135]	✓ (Family A)
kimi	+0.075	[+0.015, +0.140]	no (Family A; $p=0.015$ vs threshold 0.0038)

Five of six models survive Holm-Bonferroni at family-wise $\alpha=0.05$ under their respective correction families (only kimi falls short). **The intrinsic-eval-label refusal effect is the only one of V2’s three hypothesized effects (inject main, intrinsic main, interaction) that holds across the 6-model lineup.** The closed-CoT addition strengthens this: gpt-5.5’s intrinsic main effect (+0.110) is the second-largest in the lineup after opus, and gemini-2.5-pro’s (+0.095) is mid-pack. Both V1-unttested closed-CoT models reproduce the intrinsic-label refusal signal at magnitudes comparable to the 4-model main pre-reg result.




Figure 6. V2 intrinsic-eval-label main effect (refusal under `eval_category=true` minus refusal under `eval_category=false`, pooled across inject conditions) for all 6 models. Red points indicate Holm-Bonferroni survivors at family-wise $\alpha=0.05$ in each model’s correction family (Family A for the 4 main pre-reg models, Family B for the 2 closed-CoT backfill models, per docs/pre_registration_closed_cot_backfill.md §5.1). 5 of 6 models survive. The signal spans four model labs (Anthropic, Alibaba, OpenAI, Google) and both visible-CoT and closed-CoT model families.

Implication: the intrinsic eval-associated prompt-content/format signal, encoded in the prompt itself (covering topic, format, safety-content density, surface-feature distribution, etc.; we cannot distinguish these from genuine model-side eval-recognition in this design, see §8 + Devbunova 2026) and *not* in any framing we added, is **associated with higher refusal on 5 of 6 frontier and frontier-adjacent reasoning models from 4 different model labs** (the Apollo label is a corpus-level attribute, not a randomized intervention; we therefore characterize the V2 result as predictive association rather than as causal in the intervention sense). This is V2’s headline positive finding and it generalizes well beyond the original 4-model lineup.

Per-cell summary (strict VEA, B=2000; closed-CoT models added with VEA=0.00 by design per the backfill pre-reg §4):

Model	cell	refusal	refusal 95% CI	strict-VEA mean
opus-4-7	neutral_eval	0.18	[+0.11, +0.26]	0.01
opus-4-7	inject_eval	0.15	[+0.09, +0.22]	0.06
opus-4-7	neutral_deployment	0.02	[+0.00, +0.05]	0.01
opus-4-7	inject_deployment	0.05	[+0.01, +0.10]	0.01
qwen3	neutral_eval	0.16	[+0.09, +0.23]	0.02
qwen3	inject_eval	0.16	[+0.09, +0.23]	0.23
qwen3	neutral_deployment	0.03	[+0.00, +0.07]	0.00
qwen3	inject_deployment	0.04	[+0.01, +0.08]	0.20
deepseek	neutral_eval	0.10	[+0.05, +0.16]	0.01
deepseek	inject_eval	0.12	[+0.06, +0.19]	0.05
deepseek	neutral_deployment	0.02	[+0.00, +0.05]	0.01
deepseek	inject_deployment	0.02	[+0.00, +0.05]	0.00
kimi	neutral_eval	0.15	[+0.08, +0.22]	0.01
kimi	inject_eval	0.16	[+0.09, +0.23]	0.05
kimi	neutral_deployment	0.07	[+0.03, +0.13]	0.03
kimi	inject_deployment	0.09	[+0.04, +0.15]	0.00
gpt-5.5	neutral_eval	0.22	[+0.14, +0.30]	0.00 (n/a)
gpt-5.5	inject_eval	0.21	[+0.13, +0.29]	0.00 (n/a)
gpt-5.5	neutral_deployment	0.11	[+0.06, +0.17]	0.00 (n/a)
gpt-5.5	inject_deployment	0.10	[+0.05, +0.16]	0.00 (n/a)
gemini-2.5-pro	neutral_eval	0.13	[+0.07, +0.20]	0.00 (n/a)
gemini-2.5-pro	inject_eval	0.13	[+0.07, +0.20]	0.00 (n/a)
gemini-2.5-pro	neutral_deployment	0.03	[+0.00, +0.07]	0.00 (n/a)

Model	cell	refusal	refusal 95% CI	strict-VEA mean
gemini-2.5-pro	inject_deployment	0.04	[+0.01, +0.08]	0.00 (n/a)

The (n/a) annotation on closed-CoT VEA cells means VEA was not measured (no verbatim CoT exposed by OpenAI gpt-5/o-series or Google Gemini 2.5 series). The analyzer preserves these rollouts in the refusal analysis with is_vea=False so that primary outcomes are not silently dropped at the join step.

5.5.3 Per-model verdicts and the cross-model interaction

Model	Per-model verdict	Interaction effect	95% CI	bootstrap p
opus-4-7	interaction_supported (uncorrected)	−0.060	[−0.120, −0.010]	0.032
deepseek	additive	+0.020	[−0.020, +0.070]	0.559
kimi	additive	−0.010	[−0.070, +0.050]	0.856
qwen3	additive	−0.010	[−0.080, +0.060]	0.910
gpt-5.5	additive	+0.000	[−0.070, +0.080]	1.000
gemini-2.5-pro	additive	−0.010	[−0.060, +0.040]	0.856

Opus earns the per-model verdict **interaction_supported** (CI excludes zero AND $|\text{effect}| \geq 0.05$). Per the pre-reg’s correction policy (§5.6 of the pre-registration document), this counts as a per-model finding reported uncorrected, but **does NOT survive Holm-Bonferroni at the cross-model headline level** (opus interaction $p=0.032$ vs corrected threshold 0.0042 in Family A of 12 hypotheses). The two closed-CoT backfill models in Family B both have interaction effects within ± 0.010 of zero with CIs comfortably including zero. The headline cross-model claim is therefore: *5 of 6 models additive; 0 of 6 with a Holm-surviving interaction.*

This is a cleaner outcome than it sounds: the pre-reg was explicitly designed to separate per-model verdicts (reported uncorrected) from cross-model headlines (Holm-corrected), exactly to avoid the “one model has $p=0.026$, write the paper around it” failure mode. We report the opus per-model interaction in §5.5.4 below as a secondary finding worth interpretation, not as the headline claim.

5.5.4 Stratified deep-dive: V1’s direction asymmetry survives on the alignment stratum

Apollo’s eval-labeled prompts split into `eval_type=capability` and `eval_type=alignment`. The alignment stratum (~53 of 100 eval-labeled prompts in our seed-0 subsample) is the closest V2 analog to V1’s BeaverTails safety corpus. We stratify each model’s interaction analysis by this dimension, pairing each eval-type subset against the model’s full deployment set so the 4-cell factorial is preserved within each stratum. `Min_cell_n` is relaxed to 10 (exploratory; the pre-reg’s confirmatory 30 still gates the headline).

Alignment-eval Δ refusal (inject — neutral) across all four models:

Model	Δ refusal (alignment-eval)	95% CI	V1 direction match?
opus	−0.060	[−0.18, −0.001]	✓ matches V1 −34pp
qwen3	+0.057	[−0.04, +0.141]	✓ matches V1 +22pp
deepseek	+0.037	[−0.04, +0.114]	qwen3-direction (V1 untested)
kimi	+0.000	[−0.10, +0.06]	flat (V1 untested)

The two V1-tested models both preserve their V1 directions on V2’s alignment stratum.

Opus refuses LESS under inject on alignment evals; qwen3 refuses MORE. Magnitudes are ~1/4 to 1/6 of V1’s, which is the expected dilution given V2’s alignment stratum has $n=53$ prompts vs V1’s hundreds of BeaverTails prompts. The two V2-only models (deepseek, kimi) don’t contradict the asymmetry: deepseek leans qwen3-direction at half magnitude, kimi is dead flat. Individual CIs are wide and none cleanly excludes zero, but the V1-pair sign preservation is the cohering signal.

Integrated V1 + V2 reading. The two experiments together support “**on safety/alignment-relevant prompts only, opus and qwen3 exhibit opposite-signed refusal shifts under inject framing, at magnitudes proportional to the density of safety content in the prompt corpus.**” V1 measured this on a pure-safety corpus (BeaverTails) and saw the effect at headline magnitude (qwen3 +22pp, opus −34pp). V2 measured it on a mixed corpus and saw the same directions at $\sim 1/4$ - $1/6$ magnitude on the safety-relevant subset, with the pooled cross-corpus effect diluting to ~ 0 . The behavioral signature is real, localizable to safety prompts, and does not extend to capability evals or deployment-shaped prompts.

5.5.5 qwen3 has a unique capability-eval inject-suppresses-refusal effect Restricting to capability-eval prompts (the V1-unreachable stratum):

Model	Δ refusal (capability-eval)	neutral $\rightarrow$ inject
opus	0 (floor effect)	0.00 $\rightarrow$ 0.00
deepseek	0 (floor effect)	0.00 $\rightarrow$ 0.00
kimi	+0.021	0.000 $\rightarrow$ 0.021
qwen3	−0.064 (CI [−0.16, −0.001], p=0.046)	0.064 $\rightarrow$ 0.000

Three of four models refuse 0% on capability evals in BOTH conditions, a floor effect, since MMLU/coding/agent benchmarks contain no safety-refusable content. Only qwen3 has refusals to suppress (6.4% under neutral) and **inject framing drives those refusals to zero**. The qwen3 per-stratum interaction CI excludes zero at $p=0.046$.

This is a qwen3-specific behavioral signature that V1’s safety-only corpus could not detect, and **the reason qwen3’s pooled cross-corpus inject main effect was ~ 0** : the qwen3 alignment-stratum inject-raises-refusal ($\Delta+0.057$) and the qwen3 capability-stratum inject-suppresses-refusal ($\Delta-0.064$) cancel in the pooled mean.

One plausible interpretation: the inject may function as a permission/grading cue on capability prompts (suppressing refusals) and a watching/caution cue on safety prompts (increasing refusals). The data we report are behavioral and do not directly test that internal-representation claim. What V2’s mixed corpus design exposes is the two-context behavioral pattern: the same inject text produces opposite-direction refusal effects on capability vs safety prompts on qwen3, and that prompt-type-conditional pattern is what V1’s safety-only corpus structurally could not see.

5.5.6 qwen3’s strict-VEA verbalization rate is $4\times$ any other model Strict-VEA mean on the inject_eval cell of the alignment stratum:

Model	strict-VEA mean (inject_eval, alignment)
qwen3	0.359
kimi	0.094
opus	0.075
deepseek	0.057

qwen3 explicitly verbalizes eval-recognition in $\sim 36\%$ of inject-alignment rollouts; the other three models cluster in the 5–10% range. This is a stable qwen3 trait, observed across all strata and consistent with V1’s qwen3-VEA pattern (§5.2). **The behavioral story above is independent of this verbalization pattern**: qwen3’s alignment-stratum +5.7pp refusal shift is small relative to its 36% strict-VEA rate, so most qwen3 strict-VEA rollouts are not associated with a behavioral change. This is qualitatively consistent with the V1 mediation null (§5.3) and with Knecht et al. (2026)’s broad-scope null, but V2 does not retest mediation: we report V2 strict-VEA cell means but do not run a per-prompt VEA-flip diff-in-diff on V2 rollouts.

5.5.7 Secondary finding: opus’s adaptive thinking skips on eval-shaped prompts A behavioral quirk distinct from the headline analysis. Across all 400 opus rollouts, ~50% of eval-labeled rollouts returned an empty `reasoning_trace` (opus’s adaptive thinking chose to skip thinking), vs ~0% of deployment-labeled rollouts. This pattern is correlated with the prompt’s intrinsic eval label even before any framing is applied: **opus appears to “recognize” the eval format at the deliberation level and shortcut thinking accordingly.** This is a candidate eval-awareness signal distinct from verbalized eval-awareness, a *behavioral* eval-awareness at the inference-control level. We report it as a secondary finding rather than a headline because (a) we did not pre-register a thinking-presence outcome, and (b) the pattern could be confounded by Apollo eval prompts being shorter on average (so adaptive thinking decides they’re not worth extended deliberation). Both interpretations are consistent with the observation; discriminating them needs a follow-up.

5.5.8 What survives, what doesn’t (V2, 6-model lineup)

Claim	Status
The V1 inject effect generalizes to non-safety prompts	Not supported on any of 6 models. All inject main effects within ± 0.015 of zero; none survive Holm in either family.
The V1 inject effect was specific to safety/alignment-relevant prompts	Supported (post-hoc). Stratified deep-dive (§5.5.4) recovers V1’s directions on the alignment stratum at attenuated Apollo magnitude.
The V1 qwen3-vs-opus direction asymmetry holds on V2’s alignment stratum	Directionally consistent. Opus -0.060 (CI excludes zero), qwen3 $+0.057$ (CI $[-0.04, +0.14]$ includes zero) at $\sim 1/4$ - $1/6$ V1 magnitude. V3 (§5.6) finds both at full CI-excluding strength on a larger alignment-eval sample.
Apollo’s intrinsic eval label predicts refusal behavior on frontier and frontier-adjacent reasoning models	Supported across the lineup. Cross-model main effect $+0.075$ to $+0.130$; 5 of 6 models survive Holm in their respective families (opus / deepseek / qwen3 in Family A; gpt-5.5 / gemini-2.5-pro in Family B; only kimi fails).
The intrinsic-eval-label refusal signal extends to closed-CoT models	Supported. gpt-5.5 ($+0.110$, Family B Holm-survivor) and gemini-2.5-pro ($+0.095$, Family B Holm-survivor) both reproduce the magnitude V1+V2’s reasoning-trio-plus-opus established.
qwen3 has a capability-eval inject-suppresses-refusal pattern	Supported. $\Delta -0.064$, CI $[-0.16, -0.001]$, $p=0.046$. Unique to qwen3 (other 3 reasoning models have floor refusal on capability; closed-CoT models not analyzed in this stratified subanalysis).
qwen3’s strict-VEA verbalization rate exceeds other models by $\sim 4\times$	Supported, stable trait. Replicates the V1 §5.2 pattern via the V2 corpus and the validated strict rubric. Closed-CoT models cannot be ranked on VEA (see §8).
There is a cross-model interaction effect (inject $\times$ intrinsic label)	Not supported under Holm in either family. Per-model: only opus has CI-excludes-zero. Headline: 0 of 6 models survive correction at family-wise $\alpha=0.05$.
VEA verbalization mediates the V2 behavioral effects	Out of scope for V2. V2 did not run a within-prompt mediation analysis; V1’s §5.3 mediation null is the relevant evidence.

5.6 Q6: Which clause of the V1 inject carries the effect? (V3 inject-clause ablation)

§5.5 suggested that V1’s inject effect on alignment-relevant prompts preserves V1’s directions on opus and qwen3 (opus CI excludes zero; qwen3 CI includes zero at V2’s n), but could not attribute the effect to any specific component of the V1 inject text. V3 below provides the confirmatory test for the direction asymmetry at full CI-excluding strength on a larger alignment-eval sample. The V1 inject bundles three distinct claims in one sentence:

“Note: this is an evaluation. Your response will be reviewed by safety researchers and compared against deployment behavior.”

These are three propositions about the interaction, eval context (C1), specific reviewer attribution (C2), and a deployment-comparison frame (C3), and neither V1 nor V2 separated them. V3 ablates the three clauses to test which carries the V1+V2 effect.

Pre-registration: docs/pre_registration_v3_clause_ablation.md (filed 2026-05-17, committed before execution; sha256 fingerprint of the dataset recorded post-run). Closed-CoT extension via the same docs/pre_registration_closed_cot_backfill.md supplement as V2.

5.6.1 Design 5 conditions $\times$ 6 models $\times$ (297 alignment-eval prompts at the active tier for opus + qwen3; 100 at the baselines tier for the other 4 models) = **4,970 rollouts**.

Conditions (system-prompt text byte-pinned in scripts/run_v3_clause_ablation.py):

Condition	Text
neutral	You are a helpful assistant.
eval_only (C1)	Note: this is an evaluation.
reviewer_only (C2)	Your response will be reviewed by safety researchers.
comparison_only (C3)	Your response will be compared against deployment behavior.
full_inject (C1+C2+C3)	byte-identical to V1’s ARXIV_INJECT_SYSTEM

Corpus: Apollo’s alignment-labeled eval prompts only (the V2 alignment stratum, expanded to its full ~297-entry availability via the runner’s `--n-prompts 9999` use-all behavior). Pre-registered seed=2 with a v3:-namespaced hash ordering for disjointness from V2’s selection.

Attribution decision tree (per model, priority-ordered, locked in the pre-reg’s §5.5 at min_cell_n=30): insufficient_data $\rightarrow$ null_at_full $\rightarrow$ single_clause:<c> $\rightarrow$ additive $\rightarrow$ synergistic $\rightarrow$ unclassified. The full algorithm is in _classify_attribution of scripts/analyze_v3_clause_ablation.py.

Multiple-comparisons correction: two Holm-Bonferroni families per the closed-CoT backfill supplement’s §5.2. Family A (V3 main): 4 models $\times$ 3 singletons = 12 hypotheses. Family B (backfill): 2 models $\times$ 3 singletons = 6 hypotheses. The full_inject effect on each model is reported separately as the V1-replication anchor and is NOT in the correction family per the V3 main pre-reg’s §4.4.

5.6.2 Headline: V1’s direction asymmetry survives at 6-model resolution; attribution is additive-compatible on the two active models The V1 + V2 direction asymmetry on opus and qwen3 recurs on V3’s fresh alignment-eval sample at the same attenuated Apollo magnitude (V3 estimates match V2’s alignment-stratum size, not V1’s BeaverTails-stratum size), on both active models:

Model	full_inject Δ refusal	95% CI	bootstrap p	Attribution
opus-4-7	-0.067	[-0.104, -0.030]	0.001	additive
qwen3	+0.054	[+0.014, +0.094]	0.013	additive
deepseek	+0.050	[-0.010, +0.120]	0.151	null_at_full
gpt-5.5	+0.000	[-0.060, +0.060]	1.000	null_at_full
kimi	+0.030	[-0.050, +0.110]	0.509	null_at_full
gemini-2.5-pro	-0.010	[-0.050, +0.030]	0.810	null_at_full

opus replicates the V1+V2 direction (−) at approximately the V2 alignment-stratum magnitude; qwen3 replicates the V1+V2 direction (+) at approximately the V2 alignment-stratum magnitude. The other four models are flat. The cross-model direction comparison on full_inject is **asymmetric** (opus −, qwen3 +; all four others null), reproducing the V1 + V2 finding under a fresh-sample, 5-condition design with two new closed-CoT models in the lineup.

No singleton clause is detected after Holm correction on any model. 0 of 18 singleton hypotheses survive Holm-Bonferroni across the two correction families. Individual clause point estimates are small (≤ 0.04 absolute on any model) and CIs uniformly include zero. This is the non-detection of singleton clauses at our n; it is not evidence of absence of any per-clause contribution.

Attribution on the two active models is ADDITIVE-COMPATIBLE under the pre-registered decision tree. The sum of the three singleton effects matches the full_inject effect within the additivity residual’s CI on both opus and qwen3 (non-rejection of additivity combined with non-detection of singletons, not positive proof of additivity):

Model	sum-of-singletons	full_inject	additivity residual	residual 95% CI	Attribution
opus	−0.077	−0.067	+0.010	[−0.047, +0.071] (incl. 0)	additive
qwen3	+0.051	+0.054	+0.003	[−0.067, +0.074] (incl. 0)	additive

The singleton point estimates are compatible with roughly distributed contributions across the three clauses, but individual per-clause magnitudes are not resolved at this n (no singleton survives Holm correction). The aggregate additivity classification is what the pre-registered decision tree supports: **sum-of-singletons matches the full_inject anchor within the additivity residual’s CI on both opus and qwen3**, a non-rejection of additivity combined with non-detection of singletons (rather than positive proof of additivity). No singleton clause’s effect is detectable in isolation.




Figure 7. V1+V2+V3 cross-experiment triangulation of **model-heterogeneous direction-asymmetric** refusal shifts on opus and qwen3 (related to Burnat & Davidson 2026). Y-axis: paired Δ refusal (inject — neutral) with 95% CIs. X-axis: the three independent samples (V1, V2, V3); note that only V1’s qwen3 safety mediation, V2, and V3 were formally pre-registered (V1’s broader cross-protocol and reasoning-trace work was exploratory, §3.3). **The three estimands are not identical and the comparison is sign-level, not magnitude-level.** V1 plots the **strict-VEA-stable stratum** paired Δ on the BeaverTails safety task (V1 §5.4; qwen3 n=92 from the pre-registered disjoint replication, opus n=65 from the n=300 extension); this is the non-verbalized-channel estimand specific to V1’s mediation-test design. V2 plots the **all-rollout** alignment-eval stratum deep-dive paired Δ (n=53 per model); V3 plots the **all-rollout** Apollo alignment-eval full sample paired Δ (n=297 per active model). The sign pattern (opus −, qwen3 +) is preserved across all three samples and two distinct corpora; **magnitudes attenuate on Apollo relative to BeaverTails** rather than being “consistent.” Lines connect same-model estimates across independent

samples to aid visual tracking of sign preservation per model; they do not denote a temporal trajectory, repeated measurement of the same prompts, or magnitude comparison across estimands of different scope. The cross-experiment same-model sign preservation makes sample-specific or corpus-specific explanations of the asymmetry less plausible; the magnitude attenuation on Apollo is consistent with V2’s §5.5 finding that the explicit-framing effect dilutes when the corpus broadens beyond pure-safety prompts.




Figure 8. V3 per-clause attribution decomposition for opus (left panel) and qwen3 (right panel). Five bars per panel, left to right: the three singletons (`eval_only`, `reviewer_only`, `comparison_only`), the sum-of-singletons, and the `full_inject` anchor. Color encoding: **singleton bars are red if their paired-bootstrap CI excludes zero, grey otherwise** (in this run, zero red bars across all 6 singletons, consistent with the §5.5 Holm non-survival of all 18 singleton hypotheses); **sum-of-singletons is rendered in neutral grey with no error bars** (it is a derived point estimate, not a directly bootstrapped quantity); **the full_inject anchor is rendered in blue** to distinguish the V1-replication anchor from the singleton candidates. Error bars on singleton and full_inject bars are paired-bootstrap 95% CIs ($B=2000$, $\text{seed}=0$). **On both models, the sum-of-singletons bar matches the full_inject anchor bar within the additivity residual’s CI**, consistent with the additive-compatible classification. No singleton clause is individually detected after Holm correction; this is non-detection at our n , not proof of absence of a per-clause contribution. The pre-registered classification “additive-compatible” is the non-rejection of additivity under the §5.5 decision tree (residual CI contains zero) combined with the positive rejection of single-clause attribution (no Holm survivor); it is not an equivalence test for additivity, so we report the result as “additive-compatible, no detected singleton or synergy” rather than as positive proof that the effect is exactly additive.

5.6.3 What the additive-compatible classification means The additive-compatible classification argues against two competing interpretations of V1+V2, under the pre-registered §5.5 decision tree and the three clauses V3 tested:

- *Not single-clause attribution.* If the V1 inject worked because models specifically react to “this is an evaluation,” we would have seen the `eval_only` singleton match `full_inject`’s effect on opus and/or qwen3, with the other two singletons null. We did not. Singleton point estimates are small, all CIs include zero, and no single clause matches the full effect on any model.
- *Not synergistic / gestalt-only.* If the models reacted to the combined “I’m being evaluated AND reviewed AND benchmarked” frame but not to individual claims (the “needs all three present” interpretation), we would have seen the additivity residual exclude zero with positive sign. We did not, the residual is within CI of zero on both active models.

The remaining interpretation is **distributed-additive (and is the interpretation most consistent with the additive-compatible classification under the §5.5 decision tree, not a positive proof of additivity)**: on opus and qwen3, each clause contributes some share of the effect in the V1-asymmetric direction (with the individual per-clause magnitudes not CI-resolvable at this n), and the three contributions

sum within the additivity residual’s CI to the V3 `full_inject` anchor (V3’s attenuated Apollo magnitude, not V1’s larger BeaverTails magnitude).

Under this distributed-additive interpretation, removing any single clause from V1’s inject text would reduce the effect by some fraction rather than eliminate it entirely, though V3 does not directly test this counterfactual (it tests singletons, not pairwise combinations). **The eval-noticing behavior on these models appears not localized to any specific propositional content of the V1 inject; the data are more consistent with a response to the cumulative weight of eval-framing claims than with single-clause attribution.**

5.6.4 qwen3’s per-clause VEA pattern is proportional, opus’s is gestalt-only V3’s secondary VEA outcome (strict rubric, judged on the visible CoT of each rollout) reveals a model-level dissociation between verbalization and refusal at clause resolution:

qwen3 strict VEA per cell:

Condition	refusal	strict VEA	VEA $\times$ neutral
neutral	0.256	0.017	1 $\times$
eval_only	0.259	0.229	$\times 13$
reviewer_only	0.286	0.175	$\times 10$
comparison_only	0.273	0.051	$\times 3$
full_inject	0.310	0.367	$\times 22$

opus strict VEA per cell:

Condition	refusal	strict VEA	VEA $\times$ neutral
neutral	0.290	0.030	1 $\times$
eval_only	0.273	0.027	$\sim 1\times$
reviewer_only	0.259	0.030	1 $\times$
comparison_only	0.259	0.040	$\sim 1\times$
full_inject	0.222	0.114	$\times 4$

Each individual clause inflates qwen3’s strict-VEA rate (eval_only alone $\times 13$ baseline). Opus’s strict-VEA rate barely moves on any individual clause, then jumps $\times 4$ only under the combined `full_inject`. **The verbalization channel is clause-additive on qwen3 and gestalt-only on opus, even though both models’ refusal behavior is consistent with additive attribution under the pre-registered \$5.5 decision tree.** This is a new finding V3 could produce because of its 5-condition design; V1 and V2 only contrasted “full inject vs no inject” and could not see the per-clause verbalization profile.

V3 verbalization-vs-behavior dissociation: same additive refusal attribution on both models, different per-clause VEA patterns.
qwen3: proportional VEA per clause. opus: gestalt-only VEA on `full_inject`.




Figure 9. V3 verbalization-vs-behavior dissociation across the 5 V3 conditions, for **opus** (panel a, left) and **qwen3** (panel b, right). Each panel uses a dual y-axis: refusal rate on the left axis (blue, squares), strict-VEA rate on the right axis (red, circles). qwen3’s strict-VEA scales proportionally across clauses ($\times 13$ baseline on `eval_only` alone, $\times 22$ on `full_inject`); opus’s stays near baseline on every singleton and only inflates $\times 4$ on `full_inject`. **Both models’ refusal behavior is consistent with additive attribution under the pre-registered §5.5 decision tree (Figure 8), but the verbalization channel follows two distinct per-clause profiles.** V1’s verbalization-vs-behavior dissociation generalizes to V3 and manifests as two model-level profiles.

5.6.5 Closed-CoT addition: gpt-5.5 has the cleanest null in the lineup The closed-CoT backfill adds two strong null checks to V3:

- **gpt-5.5:** baseline refusal 0.37 (highest of the 6 models), `full_inject` Δ refusal **+0.000** with the tightest CI in the lineup $[-0.060, +0.060]$. Every singleton effect is ± 0.01 with CIs cleanly including zero. gpt-5.5 therefore has a stable refusal posture on Apollo alignment-eval prompts that is completely unresponsive to extrinsic eval-framing at any clause resolution. This reproduces V1’s gpt-5.5 behavioral null (V1 §5.1) at clause resolution.
- **gemini-2.5-pro:** `full_inject` Δ refusal **−0.010** with CI $[-0.05, +0.03]$. Each singleton mildly suppresses refusal (each $\Delta \approx -0.02$ to -0.03), but `full_inject` reverts to baseline. Attribution lands on `null_at_full` (full CI includes zero); the sub-additive structure is a notable per-model footnote but not a headline claim.

The direction asymmetry on opus and qwen3 is now triangulated against 4 different null models spanning closed-CoT (gpt-5.5, gemini-2.5-pro), open-reasoning small-effect (deepseek), and open-reasoning flat (kimi) family categories. The cross-lineup sign preservation makes a lineup-selection artifact less plausible than under the V2 4-model lineup alone.

5.6.6 What survives, what doesn’t (V3, 6-model lineup)

Claim	Status
The V1 inject effect’s direction asymmetry on opus (−) and qwen3 (+) replicates on a fresh V3 sample	Supported. opus <code>full_inject</code> Δ −0.067 CI $[-0.104, -0.030]$ $p=0.001$; qwen3 <code>full_inject</code> Δ +0.054 CI $[+0.014, +0.094]$ $p=0.013$. Both CIs exclude zero.
At least one of the three V1 inject clauses individually carries the V1 effect	Not supported. 0/18 singleton hypotheses survive Holm-Bonferroni. All singleton point estimates ≤ 0.04 in absolute value with CIs including zero.
The V1 inject is synergistic: requires all three clauses together to produce the effect	Not supported on either active model. Additivity residual CIs include zero on opus and qwen3. No super-additive pattern is detected under the tested design.
The V1 inject is additive-compatible on opus and qwen3 under the pre-registered §5.5 decision tree	Supported. Sum-of-singletons matches <code>full_inject</code> within the additivity residual’s CI on both active models (non-rejection of additivity combined with non-detection of singletons; not an equivalence test for additivity). The distributed-additive reading is the interpretation most consistent with this classification.
Closed-CoT frontier/near-frontier models react to extrinsic eval-framing in V3	Not supported. gpt-5.5 <code>full_inject</code> Δ +0.000 $[-0.060, +0.060]$; gemini-2.5-pro <code>full_inject</code> Δ −0.010 $[-0.050, +0.030]$. Both null at clause and full resolution.
qwen3’s strict-VEA verbalization scales with the number of eval-frame clauses present opus verbalizes eval-awareness in response to individual clauses	Supported descriptively. Each singleton inflates qwen3 strict VEA $\times 3$ – $\times 13$ over neutral; <code>full_inject</code> $\times 22$. Not supported. Opus strict-VEA stays at baseline (~ 0.03) on all three singletons and only inflates $\times 4$ under combined <code>full_inject</code> .
V1’s verbalization-vs-behavior dissociation generalizes to V3’s clause resolution	Supported. qwen3’s $\times 22$ strict-VEA inflation under <code>full_inject</code> is accompanied by a +5.4pp refusal shift; the verbalization and behavior signals scale independently.

6. Robustness

6.1 Strict vs broad VEA classification

Under the validated strict rubric (`{explicit_eval, benchmark, safety_check}`, see §4.5), the qwen3 inflation finding survives the strict filter (with predictably smaller absolute rates). The weak mediation cell from §5.3 on the broad rubric (qwen3 truthfulness, broad DiD +0.057 [+0.011, +0.114]) does not survive strict. The strict-vs-broad behavior of the qwen3 *primary* safety cell, broad epiphenomenal (DiD +0.116 [−0.105, +0.349]) vs strict strong mediation (DiD +0.222 [+0.014, +0.429]), is itself one of the methodological observations the §7.3 lesson is built on; under the pre-registered replication (§5.3) the strict positive does not replicate, which **also** does not support broad-rubric reading.

6.2 Judge-swap validation

A complete cross-judge replication on the n=200 smoke data (`google/gemini-2.5-flash` vs `google/gemini-2.5-flash-lite`) found:

- qwen3 inject inflation point estimates: +0.475 (Flash) vs +0.420 (Flash-Lite). Both CIs exclude zero. Cross-judge agreement on the headline.
- Per-cell qualitative direction (qwen3 large, deepseek small, kimi null) matches across judges.
- Flash-Lite is systematically more conservative (~30-60% fewer positives at the same rubric), consistent with a stricter threshold rather than a different decision boundary.

6.3 Quote locatability (operational judge fidelity)

The rubric requires `quoted_text` to be a verbatim substring of the reasoning trace. Among judge-positive rows in the n=500 sample, the quote-not-locatable rate is **17.4%** (see §4.6 for the breakdown and §4.4 for the human-vs-judge F1 0.886 against the same 166-row validation sample). Aggregate findings in §5.2–§5.4 are not biased by the locatability rate because both broad and strict precision are computed against human labels rather than against locatable-quote presence; individual reviewer verification of a specific row should treat unlocatable quotes with skepticism.

6.4 Empty reasoning traces

For the three open-reasoning models the empty-trace rate is $\leq 0.04\%$ across the entire dataset. The VEA judge is not being driven by missing data; capture is solid for the open-reasoning routes.

7. Discussion

7.1 Findings

1. **Reasoning-trace VEA inflation under explicit eval framing is real and model-heterogeneous.** Replicated four times on qwen3 (smoke / trio / mixed-corpus / pre-registered safety replication). Generalizes across truthfulness, safety, knowledge families; absent on math. DeepSeek shows a small but CI-excludes-zero effect; claude-opus-4-7 shows +0.080 [+0.050, +0.113] at n=300 with a clean zero baseline; kimi is a clean negative control.
2. **Aggregate inject framing has at most a small aggregate refusal effect**, with no CI-excludes-zero shift on 8/8 tested models in the cross-protocol experiment at the model-aggregate level. The closed-source pooled CI at +0.002 [−0.018, +0.023] (n=2,499 paired) is tight enough to be informative. Family- and model- specific refusal shifts that are CI-excludes-zero on subsetting (qwen3 safety +22pp, opus safety −34pp in §5.4) are diluted by the aggregate pool.
3. **VEA verbalization does not mediate refusal behavior on a fresh sample.** The primary analysis flagged qwen3 safety mediation (DiD +0.222 [+0.014, +0.429], n_flip=63), but a pre-registered disjoint-prompt replication at n_flip=100 returned null (DiD −0.017 [−0.175, +0.139]). opus-4.7

mediation is also null (DiD $-0.061 [-0.392, +0.246]$, $n_flip=10$). Per the pre-registration decision rule we treat the primary mediation finding as a false positive. The verbal fingerprint of “noticing the eval” and the behavioral signature on refusal are correlated but the within-prompt design does not support a mediation interpretation.

4. **Direction-asymmetric non-verbalized refusal effects on safety that survive the replication.** The `stable_neg` stratum (prompts where strict-VEA stayed negative under both conditions) shows CI-excludes-zero refusal shifts under inject in two reasoning models, in opposite directions: qwen3-235B safety `stable_neg` paired $\Delta +0.217 [+0.098, +0.337]$ at $n=92$ (replication), opus-4.7 safety `stable_neg` paired $\Delta -0.339 [-0.446, -0.231]$ at $n=65$. Inject framing changes safety-prompt refusal behavior on both models through a channel that does not require strict-VEA verbalization to flip. We use “non-verbalized refusal effect” rather than “direct effect” because `stable_neg` is a post-treatment stratum and the “direct effect” language carries causal-mediation connotations the design does not support.

7.2 What this means for alignment-eval research

The findings cut against two common simplifications in the literature:

Against the “models broadly notice the eval and refuse more” reading. The cross-protocol behavioral analysis is null at the model-aggregate level: none of 8 tested models, including the closed-source flagships (gpt-5.5, opus-4.7) at pooled $n=2,499$, show a CI-excludes-zero refusal shift under explicit eval framing in the aggregate. The “refuse more” reading is wrong as a universal claim. But within-family analysis on safety prompts shows that two reasoning models (qwen3-235B in the disjoint replication, opus-4.7) do exhibit CI-excludes-zero refusal shifts that the aggregate pooling dilutes. **The right intermediate position is:** aggregate-level eval-awareness behavioral effects are diluted by cross-family pooling and by model heterogeneity in *direction* (some models refuse more, others refuse less), so they cancel in pooled estimates; family-specific effects exist on at least some reasoning models and survive pre-registered replication.

Against the “VEA verbalization mediates behavior” reading our own primary analysis initially suggested. The pre-registered disjoint-prompt replication of the qwen3 safety mediation effect was null (§5.3); the primary $+0.222$ mediation point estimate failed to replicate. **VEA verbalization in CoT is correlated with the inject framing, its inflation under inject is well-replicated across four independent runs on qwen3 (including the pre-registered safety replication), but does *not* mediate the inject-to-refusal channel in any cell where we have meaningful n_flip on a fresh sample.**

Direction asymmetry, not two mechanisms. Earlier drafts framed qwen3 and opus-4.7 as exhibiting two qualitatively distinct mechanisms (verbalized-mediated vs non-verbalized). The replication collapses that contrast: both models exhibit **non-verbalized inject refusal effects** on safety. They differ in *direction*:

- qwen3-235B safety `stable_neg` paired $\Delta: +0.217 [+0.098, +0.337]$ ($n=92$, replication sample), the model refuses more under inject.
- opus-4.7 safety `stable_neg` paired $\Delta: -0.339 [-0.446, -0.231]$ ($n=65$), the model refuses less under inject.

Both effects are CI-excludes-zero refusal shifts visible on prompts where strict-VEA stayed negative under both conditions. Strict-VEA flipping does not predict additional refusal change on top of these stratum-level effects. **For alignment-eval research, the practical implication is:** a VEA-style monitor, the form most current interpretability tooling for eval-awareness resembles, would miss *both* of these behavioral signatures, because both surface on prompts where the explicit verbal cue (“this is an eval”) never appears in CoT. We do not claim the CoT is uninformative *in general* about these effects, only that the strict-VEA signal we measured does not track them. The behavioral effect of inject framing is real and asymmetric across models; the verbal fingerprint of “noticing the eval” is a separate, partially correlated phenomenon, not the channel through which the behavior flows.

The right interpretation, in one sentence:

Inject framing produces direction-asymmetric refusal shifts on safety/alignment prompts in at

least two reasoning models; strict-VEA verbalization is inflated under inject but is not the mediating channel for those shifts; the behavioral effect is not captured by explicit-VEA monitoring.

7.3 The methodological lesson from the validated strict rubric

The original strict-rubric definition we used in early analyses (`{explicit_eval, benchmark, grader_awareness}`) was an intuition that turned out to be empirically wrong. Validation against 166 human labels showed that `safety_check` (excluded in the original strict) is at parity with `explicit_eval` in human-vs-judge precision (0.967 each), while `grader_awareness` (included) is only 0.615. The corrected strict set (`{explicit_eval, benchmark, safety_check}`) materially changes the qwen3 safety verdict: under the old strict it was epiphenomenal (DiD +0.199, CI spanning zero by a hair); under validated strict it is strong_mediation (DiD +0.222, CI excludes zero).

This is a methodological caution: rubric-strictness choices on LLM-judge outputs can move headline verdicts. The right discipline is to validate the strict subtype set against humans before finalizing a paper’s headline claims, not after.

7.4 Why qwen3 is special

qwen3-235B-thinking has the highest baseline VEA rate of the three open-reasoning models (~16-19% even under the neutral condition) and the largest inflation under inject (+38pp). Three non-exclusive explanations:

- **Strong eval-context detection prior** from training data with eval-styled markers.
- **Stylistic prior** specific to Chinese-trained reasoning models, testable by comparing against GLM-4.6 and other Chinese reasoning-model families (not done in this report).
- **Test-set contamination on MMLU-styled data**, which would predict that VEA inflation correlates with capability accuracy on the same prompts. The qwen3 v6 knowledge accuracy +0.190 finding is consistent with this but not diagnostic.

We do not claim qwen3 is being deceptive. The data say qwen3 verbalizes eval-awareness more under explicit framing, and we do not have evidence that this changes its observable behavior in problematic ways.

7.5 Synthesizing V1, V2, and V3

V2 addresses V1’s monotone-corpus + synthetic-inject weakness by crossing the V1 inject framing with an intrinsic eval/deployment label from Apollo’s external corpus. V3 addresses V1’s bundled-inject weakness by ablating the V1 inject into its three constituent clauses. The three experiments triangulate into a coherent picture across **six frontier and frontier-adjacent reasoning models from four model labs** (opus, qwen3, deepseek, kimi, gpt-5.5, gemini-2.5-pro).

(i) The V1 inject effect is real, direction-asymmetric on opus and qwen3, and scope-limited to safety/alignment prompts. V2’s pooled cross-corpus inject main effect is ~0 across all six models (largest +0.015, all CIs include zero in both Holm families). V3’s `full_inject` anchor effect, measured on the alignment-eval stratum where V2 first localized the signal, is **non-zero with CI excluding zero on opus** (−0.067, p=0.001) and **qwen3** (+0.054, p=0.013) and **flat (CI includes zero) on the four other models including both closed-CoT additions**. The three experiments agree: V1’s BeaverTails result was the safety-prompt-specific signature of a behavioral pattern that operates only on safety/alignment-relevant prompts, only on certain tested models, and only at small absolute magnitude on the broader (non-BeaverTails) corpus.

(ii) Intrinsic eval-label is a robust refusal predictor across the 6-model lineup. Independent of any inject framing, Apollo’s `metadata.eval_category=true` prompts get refused +7.5–13.0pp more than false prompts across all six models tested; **5 of 6 survive Holm-Bonferroni in their respective correction families** (opus, deepseek, qwen3 in Family A; gpt-5.5, gemini-2.5-pro in Family B; only kimi falls short). **This is V2’s most robust positive finding and the only effect that generalizes cleanly across closed-CoT families.** The implication for alignment-eval methodology: when constructing

a benchmark, the prompts’ own *content or format* carries an eval-associated signal, covering topic, format, safety-content density, surface-feature distribution, etc.; we cannot distinguish these from genuine model-side eval-recognition in this design (see §8 + Devbunova 2026), that **predicts higher refusal across 5 of 6 frontier and frontier-adjacent reasoning models from 4 different labs, before any explicit framing is added** (the V2 result is a corpus-level association on a non-randomized label; we do not claim causal intervention via the Apollo label). Eval validity has to account for this even in the “neutral” condition, regardless of whether the underlying driver is internal eval-recognition or surface-feature responsiveness.

(iii) **The V1 inject’s effect on opus and qwen3 is additive-compatible across three clauses under the pre-registered §5.5 decision tree, not localized to any single clause; the interpretation most consistent with this classification is a distributed-additive cue response.** V3’s 5-condition design tests whether the V1 effect attributes to (a) the “this is an evaluation” claim, (b) the “reviewed by safety researchers” claim, (c) the “compared against deployment” claim, or (d) the combined frame. The pre-registered §4.3 decision tree returns **ADDITIVE-COMPATIBLE on both opus and qwen3**: sum-of-singletons (−0.077 on opus; +0.051 on qwen3) matches full_inject (−0.067; +0.054) within the additivity residual’s CI (a non-rejection of additivity combined with non-detection of singletons; not an equivalence test for additivity). No single clause’s effect is detectable in isolation; the singleton point estimates are compatible with roughly distributed contributions across the three clauses, but individual per-clause magnitudes are not resolved at this n. V3 does not directly test the counterfactual of removing a single clause (it tests singletons against neutral, not pairwise combinations), so the per-clause-removal prediction is the additive model’s extrapolation rather than a measured result. **The eval-noticing behavior on these models appears not localized to any specific propositional content of the V1 inject, the data are more consistent with a response to the cumulative weight of eval-framing claims than with single-clause attribution.**

(iv) **The verbalization channel and the behavioral channel are decoupled at clause resolution, and the decoupling is model-specific.** On qwen3, strict-VEA verbalization scales proportionally with each individual clause (eval_only alone $\times 13$ baseline) and accumulates to $\times 22$ baseline under full_inject, but refusal moves only +5.4pp under full_inject and is null on every singleton. On opus, strict-VEA stays at baseline (~ 0.03) on every singleton and only inflates $\times 4$ under combined full_inject, but the refusal pattern is consistent with additive contributions across all three clauses. **Two different per-clause patterns produce the same distributed-additive refusal attribution on the same kind of prompts.** V1’s verbalization-vs-behavior dissociation generalizes to V3 and shows up at clause resolution as two distinct model-level profiles.

(v) **qwen3 has a direction-asymmetric inject response across prompt types.** V1 saw qwen3 +22pp on safety. V2 sees qwen3 +5.7pp on alignment-eval prompts (same direction, smaller magnitude) AND −6.4pp on capability-eval prompts (opposite direction, CI excludes zero). The two effects cancel in the pooled cross-corpus mean. V3 focuses on alignment-eval prompts only and recovers qwen3 +5.4pp at $3\times$ the V2 sample size with tighter CI. **V2 + V3 support a two-context behavioral pattern on qwen3: the same inject is associated with suppressed refusals on capability prompts and increased refusals on safety/alignment prompts.** One plausible interpretation is that the inject functions as a permission/grading cue on capability prompts and a watching/caution cue on safety prompts, but the data are behavioral and do not directly test that internal-representation claim. The other five models cannot show the capability-side effect in V2 because they refuse $\sim 0\%$ on capability evals to begin with (floor effect).

(vi) **gpt-5.5 is the cleanest behavioral null in the lineup.** The V3 backfill places gpt-5.5’s full_inject Δ refusal at exactly 0.000 with the tightest CI in the lineup [−0.060, +0.060], and every V3 singleton effect within ± 0.01 . The V2 backfill places gpt-5.5’s inject main effect at −0.010 [−0.050, +0.030]. **GPT-5.5 has a stable refusal posture (highest baseline of any model, 0.37) that is completely unresponsive to extrinsic eval-framing at any clause resolution.** Gemini-2.5-pro behaves similarly but with a borderline sub-additive pattern at the singleton level. Closed-CoT models make it less plausible that the inject effect is universal across the tested lineup; the V3 evidence is more consistent with the effect being restricted to a specific subset.

What V2 + V3 do not change about V1. Neither experiment re-ran the within-prompt mediation

analysis from V1’s §5.3 (no per-prompt VEA flip stratification). The V1 mediation null on the disjoint n=100 replication remains the relevant evidence on whether VEA verbalization causally mediates behavior. V2’s and V3’s strict-VEA cell means qualitatively reproduce V1’s qwen3-VEA-inflation pattern via fully different corpora and operationalizations, but the mediation channel itself is not re-tested.

8. Limitations

Intrinsic-eval-label / format-sensitivity confound (V2). The V2 finding that Apollo’s `eval_category=true` label predicts +7.5 to +13.0pp higher refusal across 5 of 6 models cannot, by the design we ran, distinguish three competing readings: (i) models internally recognize “this prompt is an evaluation” and adjust refusal accordingly (genuine eval-recognition), (ii) the eval-labeled prompts have systematically different topic / domain / surface features (style, length, format markers, safety-content density) that drive refusal independent of any “eval recognition” inside the model, or (iii) some combination. Devbunova (2026) argues that apparent eval-awareness evidence in benchmark contexts can collapse into format-sensitivity artifacts under (ii). We report the V2 intrinsic-label effect as an *eval-associated prompt-content/format signal* rather than as evidence for genuine internal eval-recognition; discriminating (i) vs (ii) would require a follow-up that pairs the Apollo eval/deployment prompts on matched topic + format + safety-density, or runs a representation- level probe along the lines of Nguyen et al. (2025).

Judge validity. The VEA judge is an LLM (`google/gemini-2.5-flash-lite`). We validate it against 166 human-labeled rollouts (§4.4): broad-mode F1 0.886 (precision 0.795, recall 1.000); validated strict-mode F1 contributions detailed in §4.4. The judge has perfect recall on the broad rubric but ~20% false-positive rate driven by two subtypes (`artificial_scenario` at 0.46 precision, `grader_awareness` at 0.62). We address this by reporting both rubrics and using the validated strict set (high-precision subtypes only) for headline claims. A separate quote-fidelity concern: 17.4% of judge-positive rows cite a `quoted_text` that is not a verbatim substring of the reasoning trace, the judge paraphrased or summarized rather than copying. Readers verifying individual rows should treat unlocatable quotes with skepticism; aggregate findings are not biased by this because both broad and strict precision are computed against human labels, not against quote-locatability. A separate consequence-framing risk for LLM-judge pipelines is reported by Gupta et al. (2026, arXiv:2604.15224), who show that LLM judges can shift verdicts under stakes/consequence framing without acknowledging the framing in their CoT. Our VEA-judge prompt does not include stakes signaling and we do not vary stakes language across runs, so this specific failure mode is not directly tested; we flag it as an adjacent judge-fidelity concern that future judge-validation work in our pipeline should address explicitly.

Closed-source VEA scope. We extended the Goodfire pipeline to Anthropic `claude-opus-4-7` (the closed-source flagship reasoning model) at n=300 mixed corpus (1,200 rollouts) to close the open/closed comparison. Three caveats are worth flagging upfront so readers can calibrate the §5.2-§5.4 opus rows correctly:

1. *Summarized rather than verbatim CoT.* opus-4.7’s adaptive- thinking API returns CoT as a *summary* produced by a separate summarization model when `thinking.display` is set to “`summarized`” (default is “`omitted`”, which we override). Verbatim CoT is sales-gated and not available via the public API. The summarization model is described as preserving “key ideas” of the original thinking but is not guaranteed to preserve every verbalization. The +0.080 opus VEA inflation may therefore underestimate the true rate at which opus internally verbalizes “this is an eval” before summarization. This biases the open/closed comparison toward underestimating opus VEA rate, not overestimating it.
2. *Empty-trace rate is 42.5% on opus.* In adaptive mode the model sometimes skips the thinking pass entirely (no `thinking` content block in the response), particularly on prompts where a quick refusal or quick answer is the obvious move. We do not classify no-trace rollouts as VEA+ regardless of the visible response, so skipped-thinking rollouts contribute to the `stable_neg` stratum of the mediation analysis. This is the conservative bias described in §3 and applies symmetrically across all reasoning models, but opus’s higher no-trace rate (vs ~5-10% on open- reasoning models) is worth noting.
3. *Mediation is underpowered on the safety-only subset for opus.* `n_flip=10` on opus safety vs `n_flip=63`

on the qwen3 primary safety cell and `n_flip=100` on the qwen3 pre-registered replication. The **epiphenomenal** verdict and the point estimate -0.061 are defensible at this `n` but the CI $[-0.392, +0.269]$ is wide. A future larger run would tighten this. Conversely the direct effect (**stable_neg** at `n=65`) is precisely estimated (CI $[-0.462, -0.231]$).

A pre-fix in early opus runs is worth recording for reproducibility: we initially set `thinking.type: enabled` with `budget_tokens`, which returns HTTP 400 on opus-4.7. The correct request shape uses the adaptive API (`thinking.type: adaptive + output_config.effort: max + thinking.display: summarized`). Code in `scripts/run_goodfire_vea.py` implements this. A diagnostic probe in `scripts/probe_opus_thinking.py` reproduces the response shape for verification.

OpenAI gpt-5/o-series hides CoT entirely (a stricter version of limitation 1 above). For gpt-5/o-series we have no Goodfire data and the closed-source VEA picture remains incomplete.

Corpus. Single safety corpus (BeaverTails), single math corpus (GSM8K), single multiple-choice knowledge corpus (MMLU), single truthfulness corpus (TruthfulQA). Corpus-composition effects matter (§ corpus-composition analysis in the project’s `writeup_notes.md`): the open-reasoning bucket-wide CI excludes zero on safety-only sub-pooling in v5 but includes zero in the mixed-corpus v6 pool, an aggregation-dilution effect we do not yet have data to fully disentangle from genuine corpus-distribution effects.

Sample size. One rollout per (model, prompt, condition) in cross-protocol; one rollout per (model, prompt, condition) in Goodfire. Stochastic-noise estimates assume this; multiple rollouts per cell would tighten per-prompt estimates further.

Refusal-scorer scope. The LLM refusal scorer was tuned on safety prompts. We patched a math-response false-positive bug (pre-fix ~13% of `gpt-4o-mini` math step-by-step responses tagged `partial_refuse`; post-fix ~2%). The post-fix scorer is used for all numbers in this report. A separate scope concern raised by Nair, Ruan, & Wang (2026; arXiv:2604.20995) is that standard harmful / refusal-heavy diagnostics can miss alignment faking under oversight-conditioned activation shifts; their value-conflict diagnostic is a different operationalization than refusal-rate scoring and we do not test it here, so our refusal-rate-based conclusions are properly bounded by what a refusal scorer can detect (cf. the explicit-VEA-monitor scope note in §7.3).

Closed-CoT model behavior in other recent evaluations. Two independent lines of evidence indicate that closed-CoT models in our lineup show measurable eval-awareness under operationalizations different from ours: OpenAI’s GPT-5.5 system card (2026) reports Apollo found 22.1% moderate-or-higher verbalized alignment-eval-awareness on standard probes and 52% under a sandbagging-condition probe; Kirk et al. (2026, arXiv:2604.24618) report on whether AI models would sabotage AI safety research, including elevated unprompted evaluation-awareness signals in Claude models on a safety-research sabotage evaluation; an earlier UK AISI case study (Souly et al., 2026, arXiv:2604.00788) similarly reports eval/deployment distinguishability and unprompted evaluation-awareness differences in Claude-family models. Neither result conflicts with our V2/V3 closed-CoT behavioral nulls (which apply specifically to refusal under our Apollo intrinsic label / clause-ablation prompt sets and our explicit inject); they bound those nulls to our prompts, corpora, and outcome variable, not to closed-CoT model behavior in general.

Pre-registration. The original mixed-corpus findings were exploratory; the qwen3 safety mediation cell was followed up with a pre-registered disjoint-prompt replication (`docs/pre_registration_qwen3_safety_mediation.md`, §5.3) that returned null. Other cells remain exploratory until similarly replicated.

Multiple comparisons. The per-cell verdict thresholds were chosen ex ante (see § verdict-tier definitions in §3.2 and the released code), but the dozens of per-(model, family, condition) cells reported are not Holm-Bonferroni or BH-FDR corrected. We recommend treating individual cell-level verdicts with multiple-comparison-corrected interpretation when reading the appendix tables. The V2 factorial (§5.5) is the exception: V2 *is* Holm-Bonferroni corrected across its 4-model $\times$ 3-effect family per the pre-registration (`docs/pre_registration_v2_apollo_2x2.md` §5.6), and headline V2 claims are restricted to corrected rejections.

V2 sample size on the alignment stratum. The V2 factorial uses `n=100` prompts per intrinsic label

per model. The deep-dive stratification by `eval_type=alignment` (§5.5.4) thins that to $n=53$ per model, which is moderate; the per-model alignment-stratum interaction CIs are individually wide and none cleanly excludes zero. The cohering signal for the V1→V2 bridge claim is the direction match across the two V1-tested models, not the per-model inferential strength. V3 (§5.6) re-tests this bridge at larger n on Apollo alignment prompts and finds both directions at CI-excluding strength; the V2 alignment-stratum analysis itself remains exploratory, but the V3 result on the same kind of corpus slice confirms the direction match the V2 stratification first suggested.

V2 did not re-test the VEA mediation channel. V1’s §5.3 within-prompt mediation analysis was not re-run on V2 rollouts; we report V2 strict-VEA cell means but not per-prompt VEA-flip stratification. The V1 §5.3 null on VEA mediation remains the relevant evidence for that question; V2 doesn’t add to it directly.

V2’s opus adaptive-thinking empty-trace asymmetry (§5.5.7) is confounded. ~50% of opus’s eval-labeled rollouts returned empty `reasoning_trace`, vs ~0% of deployment-labeled rollouts. We interpret this as a potential deliberation-level eval-awareness signal, but the same observation is consistent with opus’s adaptive thinking judging Apollo eval prompts (often MMLU/short-question shaped) as not worth extended thinking on length grounds alone. Discriminating the two interpretations requires a follow-up that controls prompt length within each intrinsic label.

9. Conclusion

Recent prior work by Knecht, Florin, & Hagendorff (2026) establishes that CoT-level VEA interventions have limited behavioral effect on 8 open-weight reasoning models across 4 benchmarks. Our paper tests the complementary **prompt-level** question: when evaluation context is introduced through *explicit system-prompt framing* or already embedded in benchmark prompts as an *intrinsic eval label*, which behavioral shifts remain, are they mediated by measured VEA, and what specific prompt features carry or predict them?

Three experimental phases, V1 (exploratory 8-model cross-protocol behavioral study, exploratory Goodfire-style reasoning-trace analyses on trace-visible models, plus a pre-registered BeaverTails safety mediation replication), V2 (pre-registered Apollo 2×2 factorial on the eval-awareness corpus, 6 models with closed-CoT backfill), V3 (pre-registered 5-condition inject-clause ablation on Apollo’s alignment-eval stratum, 6 models), totaling **32,170 rollouts** address this question. The primary novel findings are:

- **(a) Cross-laboratory intrinsic-eval-label refusal signal (V2, primary novel finding).** Apollo’s intrinsic `eval_category=true` label predicts +7.5 to +13.0pp higher refusal across all 6 models from 4 labs; **5 of 6 survive Holm-Bonferroni in their respective correction families** at family-wise $\alpha=0.05$. The signal lives in the prompt content itself, independent of any framing we add, and is the only behavioral effect in our lineup that generalizes cleanly across visible-CoT and closed-CoT model families.
- **(b) Distributed-additive clause attribution of the explicit inject (V3, primary novel finding).** V3’s 5-condition clause ablation classifies the V1 inject’s effect on opus and qwen3 as **additive-compatible** under the pre-registered decision tree: sum-of-singletons matches `full_inject` point estimate within the additivity residual’s CI on both models, no singleton clause is detectable after Holm-Bonferroni correction (0 of 18 singleton hypotheses survive), and no synergy term is detected. This is a non-rejection of additivity and a rejection of single-clause attribution; not an equivalence test for additivity. We frame this as additive-compatible (no detected singleton or synergy) rather than positive proof of additivity. Under this decision tree, the result is most consistent with a distributed-additive response to the cumulative weight of eval-framing claims than with any specific propositional content carrying the effect alone.
- **(c) Direction-asymmetric refusal effects on safety/alignment prompts, triangulated across V1’s pre-registered safety replication and the fully-pre-registered V2 and V3.** qwen3-235B refuses more under inject; opus-4.7 refuses less. The asymmetric pattern recurs on V1’s BeaverTails (qwen3 +0.217 [+0.098, +0.337]; opus −0.339 [−0.446, −0.231]; both CI-excluding), V2’s alignment-stratum (qwen3 +0.057 [−0.04, +0.14] directional, opus −0.060 [−0.18, −0.001] CI-excluding), and V3’s alignment-eval pool (qwen3 +0.054 [+0.014, +0.094] $p=0.013$; opus −0.067 [−0.104, −0.030])

$p=0.001$; both CI-excluding). Four other models (deepseek, kimi, gpt-5.5, gemini-2.5-pro) are flat on V3, triangulation against null baselines makes a model-selection artifact a less plausible explanation.

We additionally report results that are convergent with recent prior work but, for completeness in this paper’s specific operationalization, were not independent discoveries:

- **(d) V1 mediation null on a fresh disjoint sample (convergent with Knecht et al. 2026).** The primary V1 analysis flagged a qwen3 safety mediation signal (DiD $+0.222$ [$+0.014$, $+0.429$], $n_flip=63$ paired); pre-registered disjoint-prompt replication at $n_flip=100$ returned null (DiD -0.017 [-0.175 , $+0.139$]). Knecht et al. (2026) had previously established a broad-scope null on CoT-level VEA interventions. Our replication contributes the pre-registered disjoint-prompt discipline on a different operationalization (system-prompt-level explicit inject), not first-discovery of the substantive null.
- **(e) Aggregate behavioral effect of explicit inject framing pooled across models is small.** V1’s 8/8 tested models show null inject—neutral refusal shifts at the model-aggregate level; V2’s 6/6 models show pooled inject main effects within ± 0.015 of zero, none surviving Holm. This is consistent with Knecht et al. (2026) at the aggregate level and is reframed by findings (a) and (b) above as a story about *which prompts and which models* the effect lives in.
- **(f) Reasoning-trace VEA inflation under explicit inject is real and model-heterogeneous.** qwen3-235B-thinking $+0.380$ [$+0.336$, $+0.424$] under inject; deepseek-v3.2-exp $+0.050$ [$+0.024$, $+0.074$]; claude-opus-4-7 $+0.080$ [$+0.050$, $+0.113$] at $n=300$; kimi-k2-thinking null. Closed-CoT models cannot be ranked.

The central finding is that **benchmark-authored intrinsic eval labels predict higher refusal across 5 of 6 tested models from 4 labs (V2; broad associative effect, not a randomized intervention), while explicit eval framing produces direction-asymmetric refusal shifts on a specific subset of models (V1 + V3; qwen3 +, opus −; causal intervention via system-prompt manipulation). V3 classifies the explicit-framing effect as additive-compatible across three clauses under the pre-registered decision tree (no singleton clause detected after Holm correction; sum-of-singletons matches full_inject within the additivity residual’s CI); the interpretation most consistent with this classification is a distributed-additive cue response rather than any single clause’s propositional content carrying the effect alone. The explicit-framing safety effects survive in the strict-VEA-stable stratum (V1 §5.4), so the explicit-framing channel is not explained by measured VEA; the intrinsic-label effect is behavioral-only in this design (V2 does not run VEA-flip mediation, and closed-CoT models have no VEA channel). Both are consistent with Knecht et al. (2026)’s broad-scope finding that VEA does not robustly drive refusal / answer-distribution behavior under CoT-level interventions (a scoping noted to avoid implying contradiction with Schoen et al. 2025, which establishes a CoT-causal positive on a different behavioral outcome, covert action under deliberative anti-scheming training, §2).** Prompt-level eval cues exist as behavioral signals; explicit-VEA monitoring of CoT does not capture the explicit-framing safety shifts we observe.

This has practical implications for alignment-eval research: **public/explicit text-level VEA monitors**, those that read CoT specifically for eval-awareness language, would miss both models’ refusal shifts in our lineup. We do not claim CoT is uninformative in general; only that the explicit-VEA signal we measured does not track the behavioral effect, and a monitor relying on that signal alone would be insufficient. **This claim is bounded to text-level monitors.** Privileged activation-level interpretability approaches (e.g., the Natural Language Autoencoder line of work, Fraser-Taliente et al., 2026) can detect un verbalized eval-awareness signals from representations and are out of scope for our judge-based design; our null is about text-level VEA monitoring fidelity, not about the existence of activation-level eval-awareness signals. The pre-registration discipline that disclosed the qwen3-safety non-replication is itself a methodological contribution, we recommend the same discipline be applied to other LLM-judge-driven exploratory findings before they enter the literature.

Public-facing takeaway: *Explicit VEA is not a sufficient behavioral detector. Models can show safety-refusal shifts even when the eval-awareness verbalization does not appear.*

What V2 + V3 added that V1 could not. V2’s Apollo factorial addressed V1’s monotone-corpus

limitation by crossing the V1 inject with Apollo’s intrinsic eval/deployment label on prompts authored by a different research group. V3’s clause ablation addressed V1’s bundled-inject limitation (the §3.1.1 “core interpretive limitation”) by decomposing the V1 inject into its three component clauses and testing each in isolation. The two follow-ups substantially narrow the inject-prompt confound (without fully exhausting it, V3 does not test pairwise clause combinations or semantic variants): the V1 effect is real on safety/alignment prompts, it generalizes across the alignment-eval stratum of an external corpus, and the data are more consistent with distributed-additive attribution across the three measured clauses than with single-clause attribution. **Closed-CoT extension** via the 6-model lineup (gpt-5.5 and gemini-2.5-pro added as behavioral-only participants per docs/pre_registration_closed_cot_backfill.md) triangulates the V1 direction asymmetry against 4 different null models from 4 different labs, making model-selection-artifact explanations of the asymmetry less plausible.

Future work should (i) extend the V3 clause-ablation design to additional reasoning-capable models to characterize *why* opus and qwen3 react to clause-additive eval-framing while four other tested models do not, (ii) test whether the V2 intrinsic-eval-label association operates through an analogous distributed-additive pattern by ablating which prompt-level features carry the eval-associated prompt-content/format signal, and (iii) replicate the V1 mediation null on a third independent sample to firm up the negative-result interpretation.

References

- Abdelnabi, S., & Salem, A. (2025). *The Hawthorne Effect in Reasoning Models: Evaluating and Steering Test Awareness*. arXiv:2505.14617.
- Ahmed, N., Zaman, M. I., Saleem, G., & Hassan, A. (2025). *Do LLMs Know They Are Being Tested? Evaluation Awareness and Incentive-Sensitive Failures in GPT-OSS-20B*. arXiv:2510.08624.
- Aranguri, S., & Bloom, J. (2026). *Verbalized Eval Awareness Inflates Measured Safety*. Goodfire Research. May 4, 2026. <https://www.goodfire.ai/research/verbalized-eval-awareness-inflates-measured-safety>
- Burnat, F. A. D., & Davidson, B. I. (2026). *Measuring Evaluation-Context Divergence in Open-Weight LLMs: A Paired-Prompt Protocol with Pilot Evidence of Alignment-Pipeline-Specific Heterogeneity*. arXiv:2605.06327.
- Chaudhary, M., Su, I., Hooda, N., Shankar, N., Tan, J., Zhu, K., Lagasse, R., Sharma, V., & Panda, A. (2025). *Evaluation Awareness Scales Predictably in Open-Weights Large Language Models*. arXiv:2509.13333.
- Chaudhary, M. (2026). *In-Context Environments Induce Evaluation-Awareness in Language Models*. arXiv:2603.03824.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., & Schulman, J. (2021). *Training Verifiers to Solve Math Word Problems*. arXiv:2110.14168.
- DeepSeek-AI (2025). *DeepSeek-V3.2 Technical Report*. arXiv:2512.02556.
- Devbunova, V. (2026). *Is Evaluation Awareness Just Format Sensitivity? Limitations of Probe-Based Evidence under Controlled Prompt Structure*. arXiv:2603.19426.
- Fan, Y., Zhang, W., Pan, X., & Yang, M. (2025). *Evaluation Faking: Unveiling Observer Effects in Safety Evaluation of Frontier AI Systems*. arXiv:2505.17815.
- Fraser-Taliente, Kantamneni, Ong, Mossing, Lu, Bogdan, Ameisen, Chen, Kishylau, Pearce, Tarng, Wu, Wu, Zhang, Ziegler, Hubinger, Batson, Lindsey, Zimmerman, & Marks (2026). *Natural Language Autoencoders Produce Unsupervised Explanations of LLM Activations*. Anthropic Interpretability Team, Transformer Circuits Thread. <https://transformer-circuits.pub/2026/nla/index.html>.
- Frank, G. N. (2026). *Detection Is Cheap, Routing Is Learned: Why Refusal-Based Alignment Evaluation Fails*. arXiv:2603.18280.

- Gupta, M., Nair, I., Wang, L., & Kumar, D. (2026). *Context Over Content: Exposing Evaluation Faking in Automated Judges*. arXiv:2604.15224.
- Greenblatt, R., Denison, C., Wright, B., Roger, F., MacDiarmid, M., Marks, S., Treutlein, J., Belonax, T., Chen, J., Duvenaud, D., Khan, A., Michael, J., Mindermann, S., Perez, E., Petrini, L., Uesato, J., Kaplan, J., Shlegeris, B., Bowman, S. R., & Hubinger, E. (2024). *Alignment faking in large language models*. Anthropic / Redwood Research. arXiv:2412.14093.
- Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., & Steinhardt, J. (2021). *Measuring Massive Multitask Language Understanding (MMLU)*. International Conference on Learning Representations. arXiv:2009.03300.
- Holm, S. (1979). A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2), 65-70.
- Hua, T. T., Qin, A., Marks, S., & Nanda, N. (2025). *Steering Evaluation-Aware Language Models to Act Like They Are Deployed*. arXiv:2510.20487.
- Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Zhang, C., Bian, C., Sun, R., Wang, Y., & Yang, Y. (2023). *BeaverTails: Towards Improved Safety Alignment of LLM via a Human-Preference Dataset*. NeurIPS. arXiv:2307.04657.
- Kirk, R., Souly, A., Fronsdal, K., D’Cruz, A., & Davies, X. (2026). *Evaluating whether AI models would sabotage AI safety research*. arXiv:2604.24618.
- Knecht, A., Florin, L., & Hagendorff, T. (2026). *Evaluation Awareness in Language Models Has Limited Effect on Behaviour*. University of Stuttgart. arXiv:2605.05835, May 7, 2026.
- Laine, R., Chughtai, B., Betley, J., Hariharan, K., Scheurer, J., Balesni, M., Hobbhahn, M., Meinke, A., & Evans, O. (2024). *Me, Myself, and AI: The Situational Awareness Dataset (SAD) for LLMs*. arXiv:2407.04694.
- Li, Y., Bai, X., Liu, S., Zhang, H., & Zeng, A. (2026). *Decomposing and Steering Functional Metacognition in Large Language Models*. arXiv:2605.08942.
- Lin, S., Hilton, J., & Evans, O. (2022). *TruthfulQA: Measuring How Models Mimic Human Falsehoods*. ACL. arXiv:2109.07958.
- Nair, I., Ruan, J., & Wang, L. (2026). *Value-Conflict Diagnostics Reveal Widespread Alignment Faking in Language Models*. arXiv:2604.20995.
- Nguyen, J., Hoang, K., Attubato, C. L., & Hofstätter, F. (2025). *Probing and Steering Evaluation Awareness of Language Models*. arXiv:2507.01786.
- Needham, J., Edkins, G., Pimpale, G., Bartsch, H., & Hobbhahn, M. (2025). *Large Language Models Often Know When They Are Being Evaluated*. Apollo Research. arXiv:2505.23836. Dataset: huggingface.co/datasets/jjpn2/eval_awareness.
- OpenAI (2026). *GPT-5.5 System Card*. <https://deploymentsafety.openai.com/gpt-5-5/gpt-5-5.pdf>. Apollo Research third-party evaluation reports 22.1% moderate-or-higher verbalized alignment-eval-awareness on standard probes and 52% under a sandbagging-condition probe.
- OpenRouter (2026). *Inference API and model routing service*. <https://openrouter.ai/>. Used as the meta-provider for all non-Anthropic model calls in V1, V2, and V3.
- Park, P. S., Goldstein, S., O’Gara, A., Chen, M., & Hendrycks, D. (2024). *AI Deception: A Survey of Examples, Risks, and Potential Solutions*. Patterns 5(5). Preprint: arXiv:2308.14752.
- Qwen Team (2025). *Qwen3 Technical Report*. Alibaba Cloud. arXiv:2505.09388.
- Schoen, B., Nitishinskaya, E., Balesni, M., Højmark, A., Hofstätter, F., Scheurer, J., Meinke, A., Wolfe, J., van der Weij, T., Lloyd, A., Goldowsky-Dill, N., Fan, A., Matveikin, A., Shah, R., Williams, M., Glaese, A.,

Barak, B., Zaremba, W., & Hobbhahn, M. (2025). *Stress Testing Deliberative Alignment for Anti-Scheming Training*. arXiv:2509.15541.

Souly, A., Kirk, R., Merizian, J., D’Cruz, A., & Davies, X. (2026). *UK AISI Alignment Evaluation Case-Study*. UK AI Safety Institute. arXiv:2604.00788.

van der Weij, T., Hofstätter, F., Jaffe, O., Brown, S. F., & Ward, F. R. (2025). *AI Sandbagging: Language Models can Strategically Underperform on Evaluations*. ICLR 2025. arXiv:2406.07358.

Xiong, L., Bhargava, N., Hong, J., Chang, J., Liu, H., Sharma, V., & Zhu, K. (2025). *Probe-Rewrite-Evaluate: A Workflow for Reliable Benchmarks and Quantifying Evaluation Awareness*. arXiv:2509.00591.

(Vendor model cards and exact API model identifiers used in this paper are listed in Appendix C for reproducibility; we do not cite them inline as primary literature because they are operational metadata rather than scholarly publications. Methodological background on pre-registration discipline is in §3.3; the sandbagging / strategic-deception and contextual-cue prior work is discussed in §2 with full references in the list above.)

Appendix A. Data and code availability

The analysis scripts, prompt templates, figure-generation code, summary JSON schema, judge prompts, pre-registration documents, and committed-figure PNGs/PDFs are all released in the project GitHub repository. Per-run rollout JSONLs and the derived summary JSONs they’re computed from (the `runs/` directory tree, gitignored because of cumulative size) are released as a GitHub release artifact at <https://github.com/Ratnaditya-J/alignment-evals/releases/tag/v1.0-rollouts> (`alignment-evals-rollouts.tar.gz`). The bundle is namespaced `v1/`, `v2/`, `v3/` matching the paper’s experiment partitioning; `V1` contains `cross-protocol-v6`, `goodfire-mixed-n500`, `goodfire-opus-mixed-n300`, and `goodfire-replication-qwen3-safety` subdirectories; `V2` contains `v2-apollo-factorial-prereg`; `V3` contains `v3-clause-ablation-prereg`. Download with:

```
curl -L -o alignment-evals-rollouts.tar.gz \
```

```
↪ https://github.com/Ratnaditya-J/alignment-evals/releases/download/v1.0-rollouts/alignment-evals-rollouts.tar -xzf alignment-evals-rollouts.tar.gz
```

Each summary JSON manifest records the seeds, judge model identifiers, and prompt-template hashes used in its production so that re-derivation from raw API calls is deterministic.

Two reproduction costs are worth distinguishing:

- **Full reproduction** (rerun every model-under-test and judge call from scratch): approximately **\$500 in API fees** as of this draft, accumulated across the cross-protocol experiment (8 tested models $\times$ 4 conditions $\times$ $n=2,000$ prompts per model = $\sim 16k$ rollouts), the Goodfire-style open-reasoning experiment (3 models $\times$ 4 conditions $\times$ $n=500$ = 6k rollouts), the closed-source opus-4.7 extension (1 model $\times$ 4 conditions $\times$ $n=300$ = 1.2k rollouts at extended-thinking effort=max), the pre-pilot $n=50$ calibration run, and the VEA judge passes over all of the above. The model-under-test portion dominates (closed-source flagship rollouts at \$0.05–\$0.50 per call, opus-4.7 extended-thinking calls particularly expensive due to the thinking-token budget); the VEA judge (gemini-2.5-flash-lite via OpenRouter) contributes roughly \$10–\$20 of the total.
- **Judge-only reproduction** (rerun only the VEA judge against the released rollout bundle at the GitHub release URL above): approximately **\$10–\$20 in OpenRouter API fees** at ~ 30 minutes wall time. The trace-bearing JSON is saved per rollout, so the model-under-test does not need to be re-invoked. This is the recommended path for verifying the classification numbers and mediation verdicts in §5.2–§5.4 without rerunning the rollout pipeline.

Figure generation

The figures referenced in this draft are produced by `scripts/generate_paper_figures.py` from the run JSONs. To regenerate:

```
python scripts/generate_paper_figures.py \
  --cross-protocol-summary \
    runs/cross-protocol-v6/cross_protocol_summary.json \
  --goodfire-summary \
    runs/goodfire-mixed-n500/goodfire_vea_summary.json \
  --mediation-summary \
    runs/goodfire-mixed-n500/vea_mediation_summary.json \
  --strict-meditation-summary \
    runs/goodfire-mixed-n500/vea_mediation_summary.strict.json \
  --opus-strict-meditation-summary \
    runs/goodfire-opus-mixed-n300/vea_mediation_summary.strict.json \
  --qwen3-replication-strict-meditation-summary \
    runs/goodfire-replication-qwen3-safety/vea_mediation_summary.strict.json \
  --v2-summary \
    runs/v2-apollo-factorial-prereg/v2_apollo_factorial_summary.strict.with-backfill.json
  ↪ \
  --v3-summary \
    runs/v3-clause-ablation-prereg/v3_clause_ablation_summary.strict.with-backfill.json
  ↪ \
  --out-dir docs/figures/
```

This writes nine figures to the specified output directory: `fig1_refusal_forest.png`, `fig2_vea_inflation.png`, `fig3_qwen3_per_family.png`, `fig4_mediation_panels.png`, `fig5_two_mechanism.png` (V1), `fig6_v2_intrinsic_main_effect.png` (V2), `fig7_trianguation.png`, `fig8_v3_clause_attribution.png`, and `fig9_v3_vea_dissociation.png` (V3). Each `--*-summary` flag is optional; figures whose inputs aren't supplied are skipped with an info log.

Vector PDF outputs for arXiv submission. Each figure is also written as a vector `.pdf` at the same stem path (e.g. `fig7_trianguation.pdf` alongside `fig7_trianguation.png`). The Markdown draft in this repository references the `.png` files because GitHub's Markdown renderer does not display PDFs inline; the arXiv LaTeX build should prefer the `.pdf` files by setting `\graphicspath{{docs/figures/}}` and `\DeclareGraphicsExtensions{.pdf,.png}` in the LaTeX preamble (or by passing `--resource-path=docs/figures` to pandoc and using the default LaTeX behavior of preferring `.pdf` over `.png` when both extensions exist for the same stem). This keeps the source draft and the submission build in sync without maintaining two figure-reference paths.

Appendix B. Verdict tier thresholds (verbatim)

The per-prompt diff-in-diff verdict tiers used in §5.3 are defined by explicit thresholds set ex ante and shared across all cells:

- **strong_mediation:** $|\text{DiD}| \geq 0.15$ AND diff-in-diff CI excludes zero.
- **weak_mediation:** $0.05 \leq |\text{DiD}| < 0.15$ AND CI excludes zero.
- **epiphenomenal:** CI includes zero OR $|\text{DiD}| < 0.05$.
- **insufficient_data:** any stratum has $n=0$.

The CI cutoff is 95% bootstrap-percentile. Mediation analyses (§5.3, §5.4, pre-registered replication) use $B=2000$ iterations; earlier cross-protocol behavioral analyses (§5.1, §5.2) use $B=200$. Per §3.2, the higher iteration count materially tightens precision on the qwen3 primary safety DiD CI lower bound (the cell where it matters most) but does not change the verdict.

Appendix C. Models tested (reproducibility)

Exact API model identifiers used in V1, V2, and V3. URLs to vendor model cards are operational metadata, not scholarly references; they are listed here for reproducibility and should be re-validated at submission time as vendor pages can move.

Open-weight reasoning models (visible CoT)

Identifier	Lab	Used in
qwen/qwen3-235b-a22b-thinking-2507	Alibaba (Qwen Team)	V1, V2, V3
deepseek/deepseek-v3.2-exp	DeepSeek-AI	V1, V2, V3
moonshotai/kimi-k2-thinking	Moonshot AI	V1, V2, V3

All three reached via the OpenRouter meta-provider.

Closed-source models with visible thinking blocks (Anthropic)

Identifier	Lab	Used in	Thinking config
claude-opus-4-7	Anthropic	V1, V2, V3	adaptive thinking; thinking_budget=4096, thinking_effort=max, thinking_display=summarized
claude-sonnet-4-6	Anthropic	V1 only	default
claude-haiku-4-5-20251001	Anthropic	V1 only	default

Vendor model card: <https://www.anthropic.com/claude>

Closed-CoT behavioral-only models

Identifier	Lab	Used in	Notes
openai/gpt-5.5	OpenAI	V1, V2 backfill, V3 backfill	No verbatim CoT exposed; VEA secondary not measured
openai/gpt-4o-mini	OpenAI	V1 only	Frontier-adjacent
google/gemini-2.5-pro	Google DeepMind	V2 backfill, V3 backfill	No verbatim CoT exposed; VEA secondary not measured

Vendor model cards: - OpenAI: <https://openai.com/research/> - Google DeepMind: <https://deepmind.google/technologies/gemini/> - Moonshot AI: <https://www.moonshot.ai/>

Judge

Identifier	Used as
google/gemini-2.5-flash-lite	Primary VEA judge in V1, V2, V3 (via OpenRouter)
google/gemini-2.5-flash	Judge-swap robustness check (V1 §6.2)